



In-memory Object Graph Stores



Aditya Thimmaiah



Zijian Yi



Joseph Kenis



Christopher Rossbach

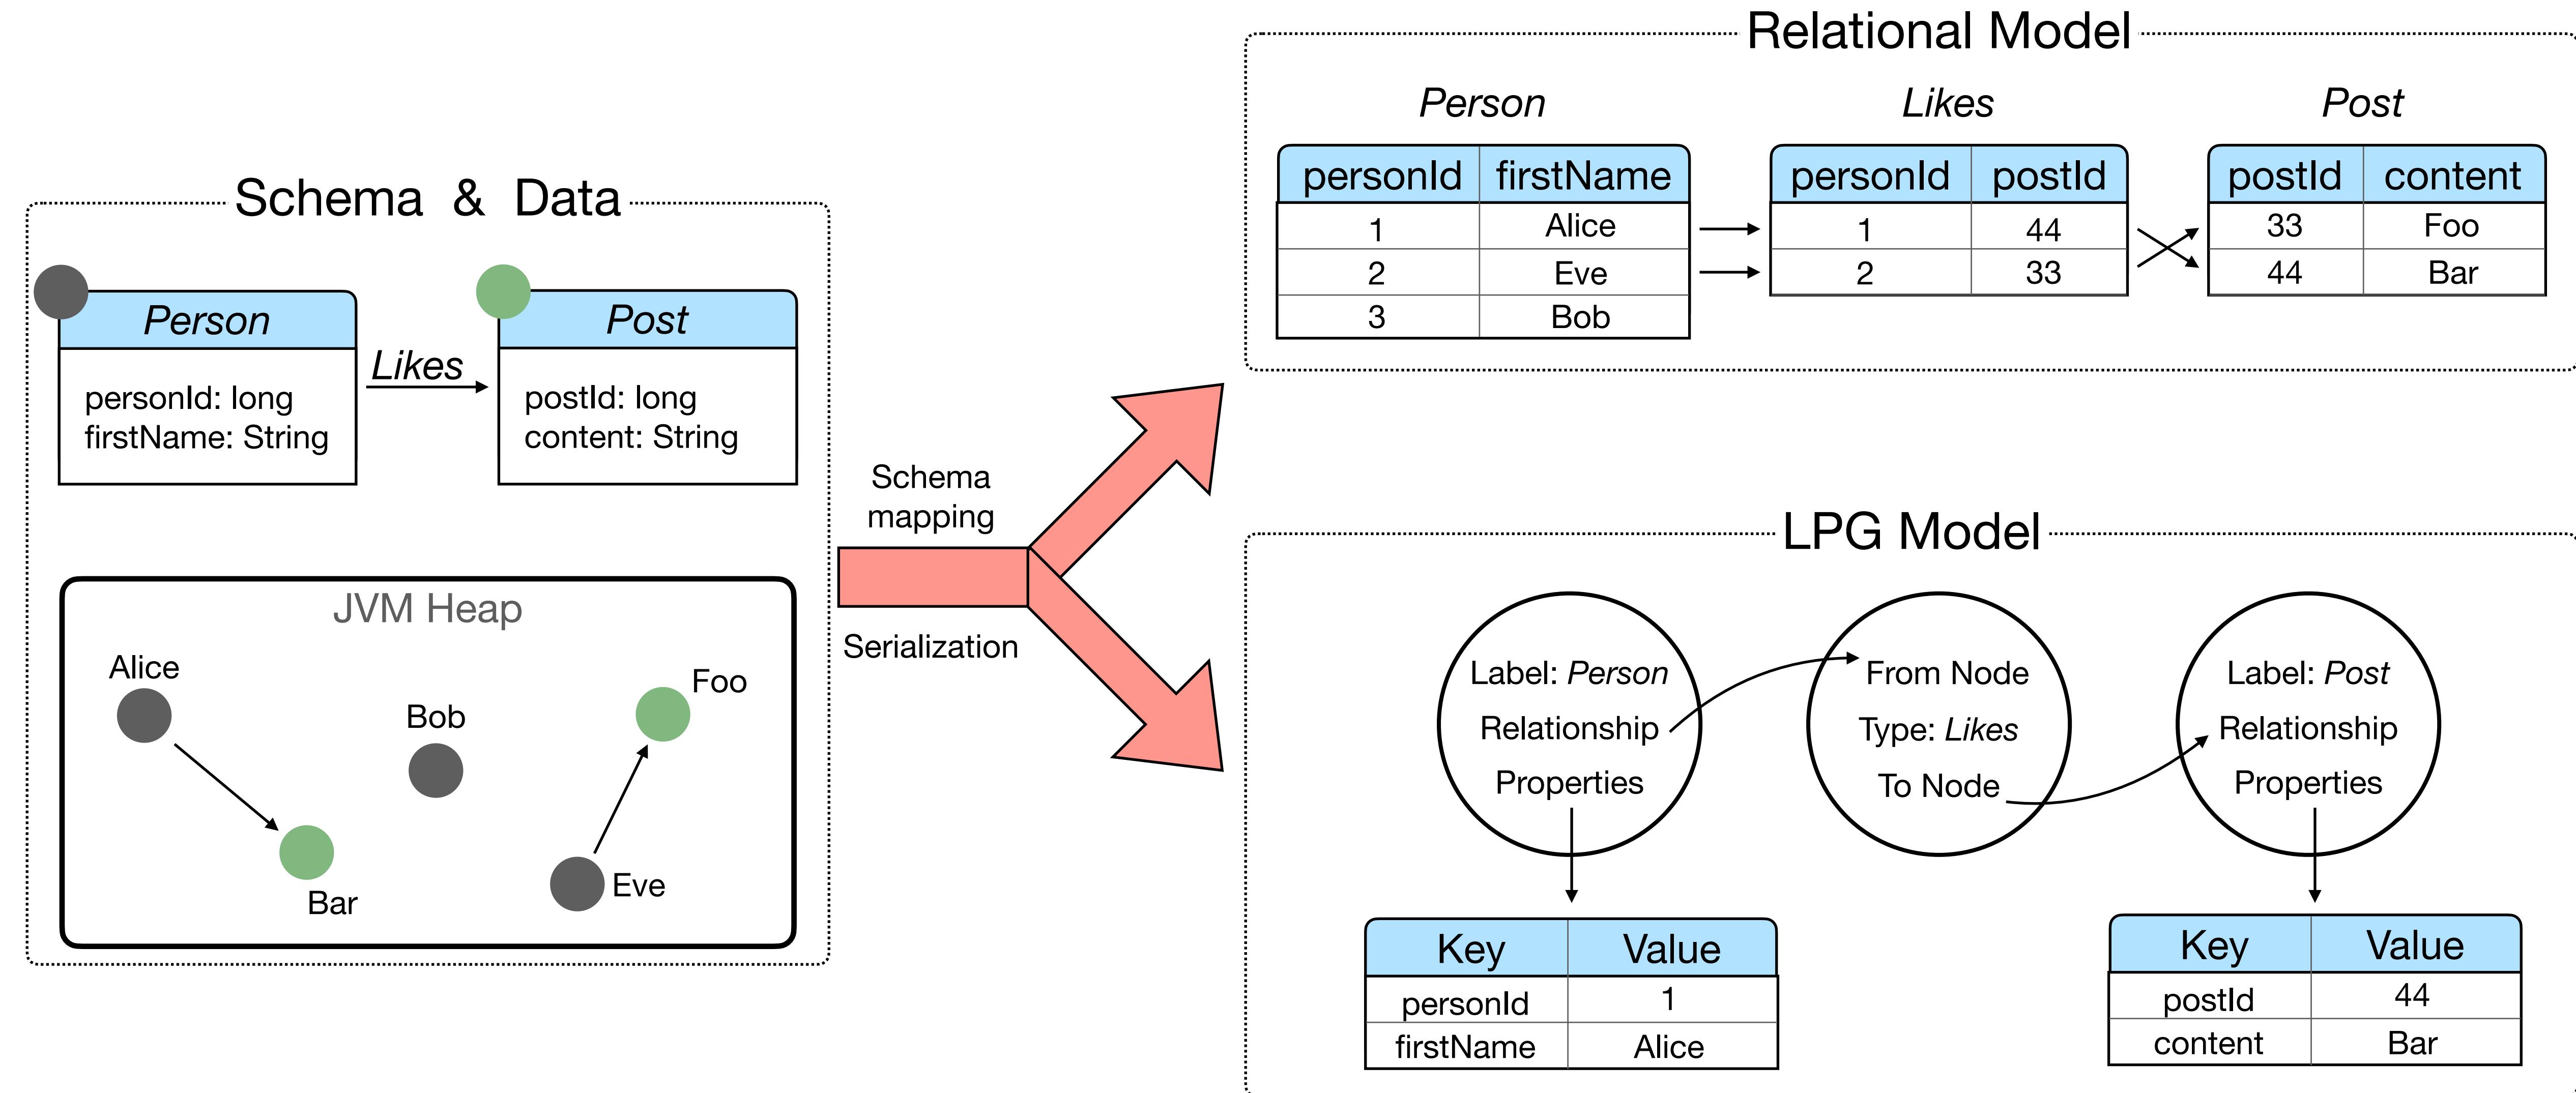


Milos Gligoric



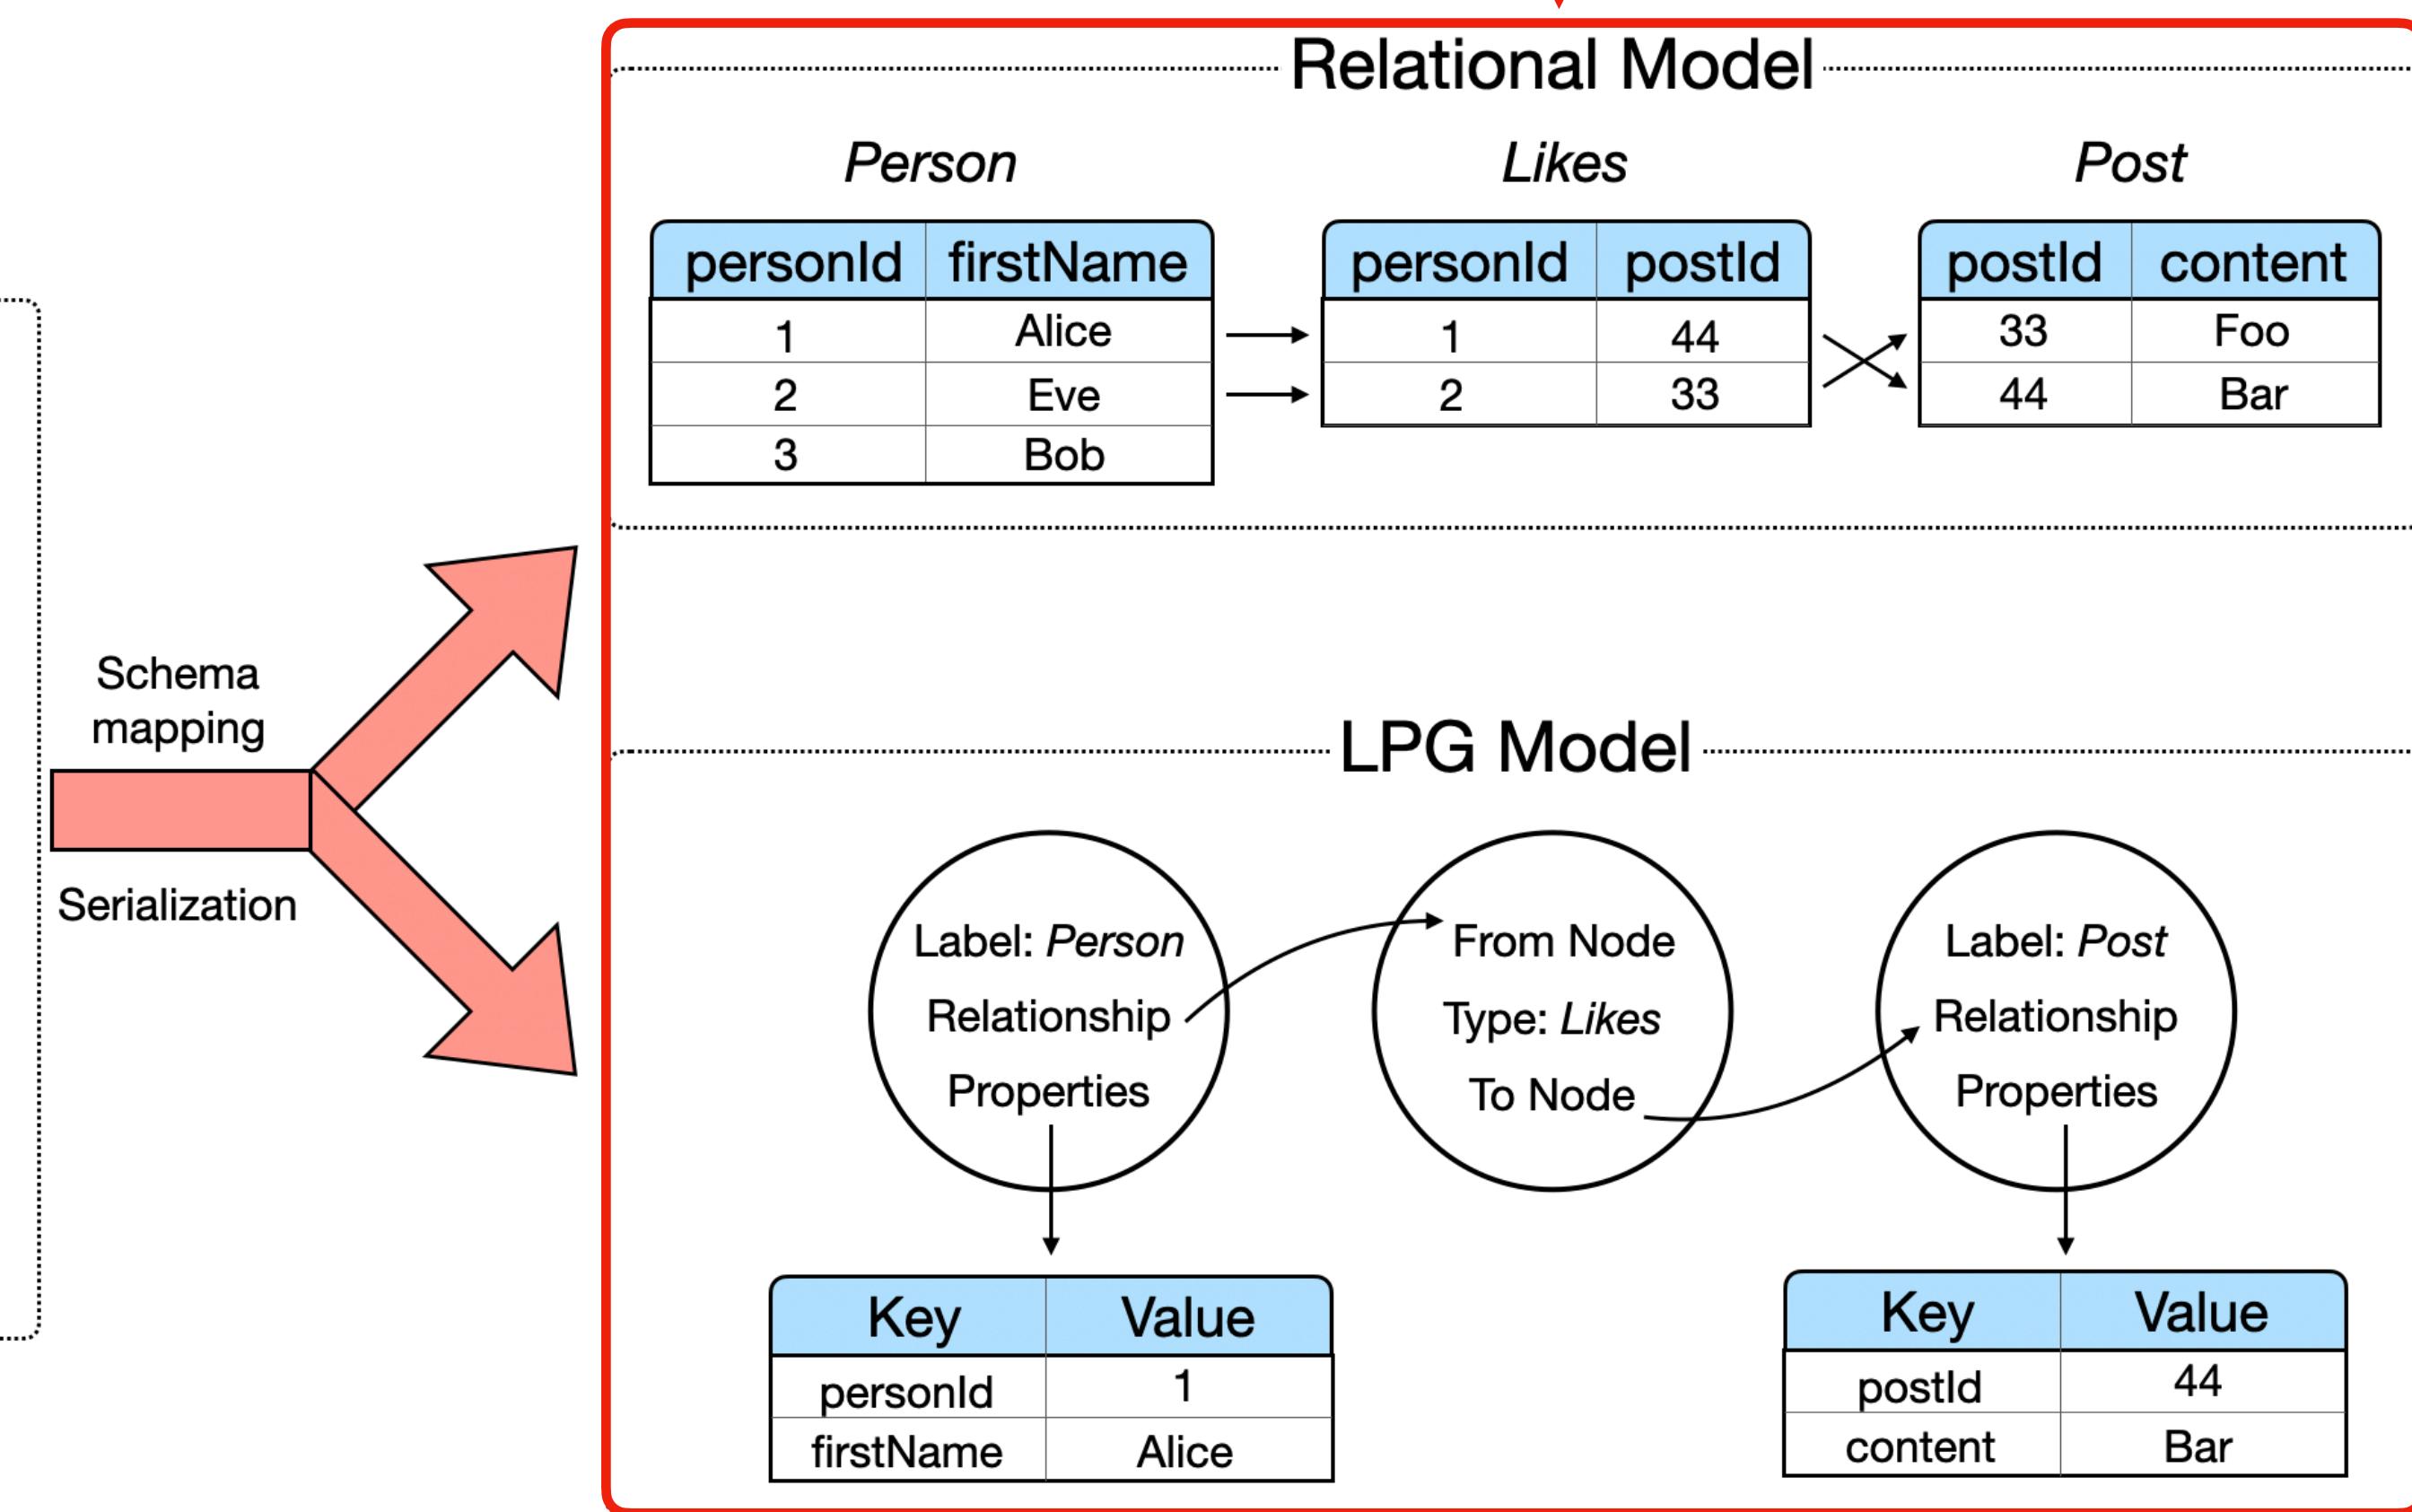
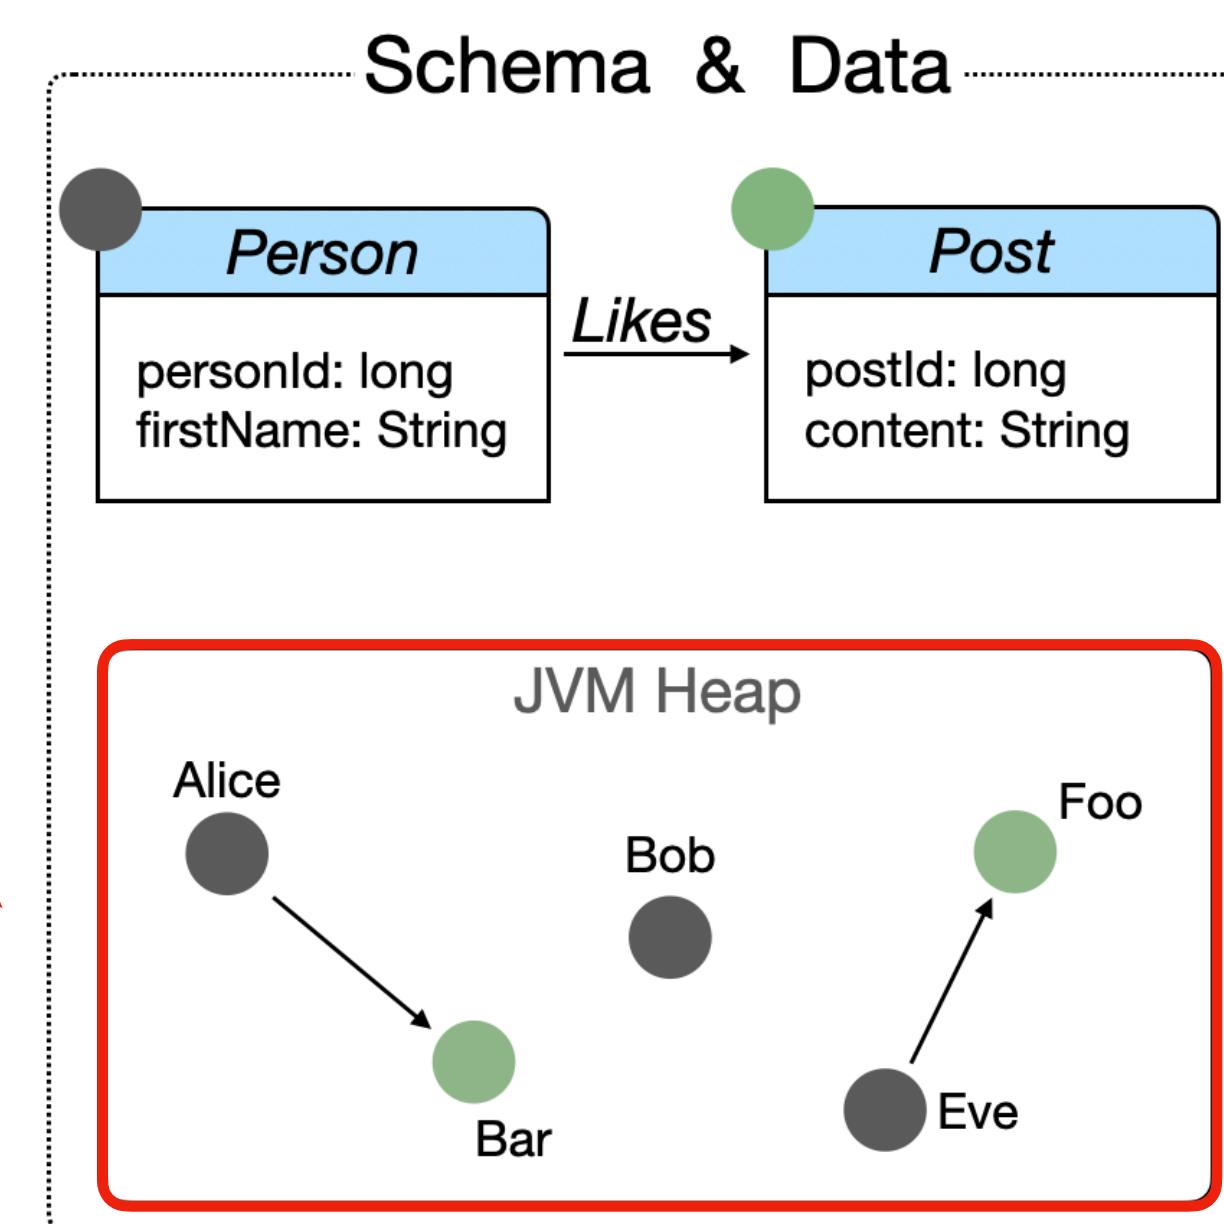
CCF-2107291
CCF-2217696
CCF-2313027
CCF-2403036

Motivation



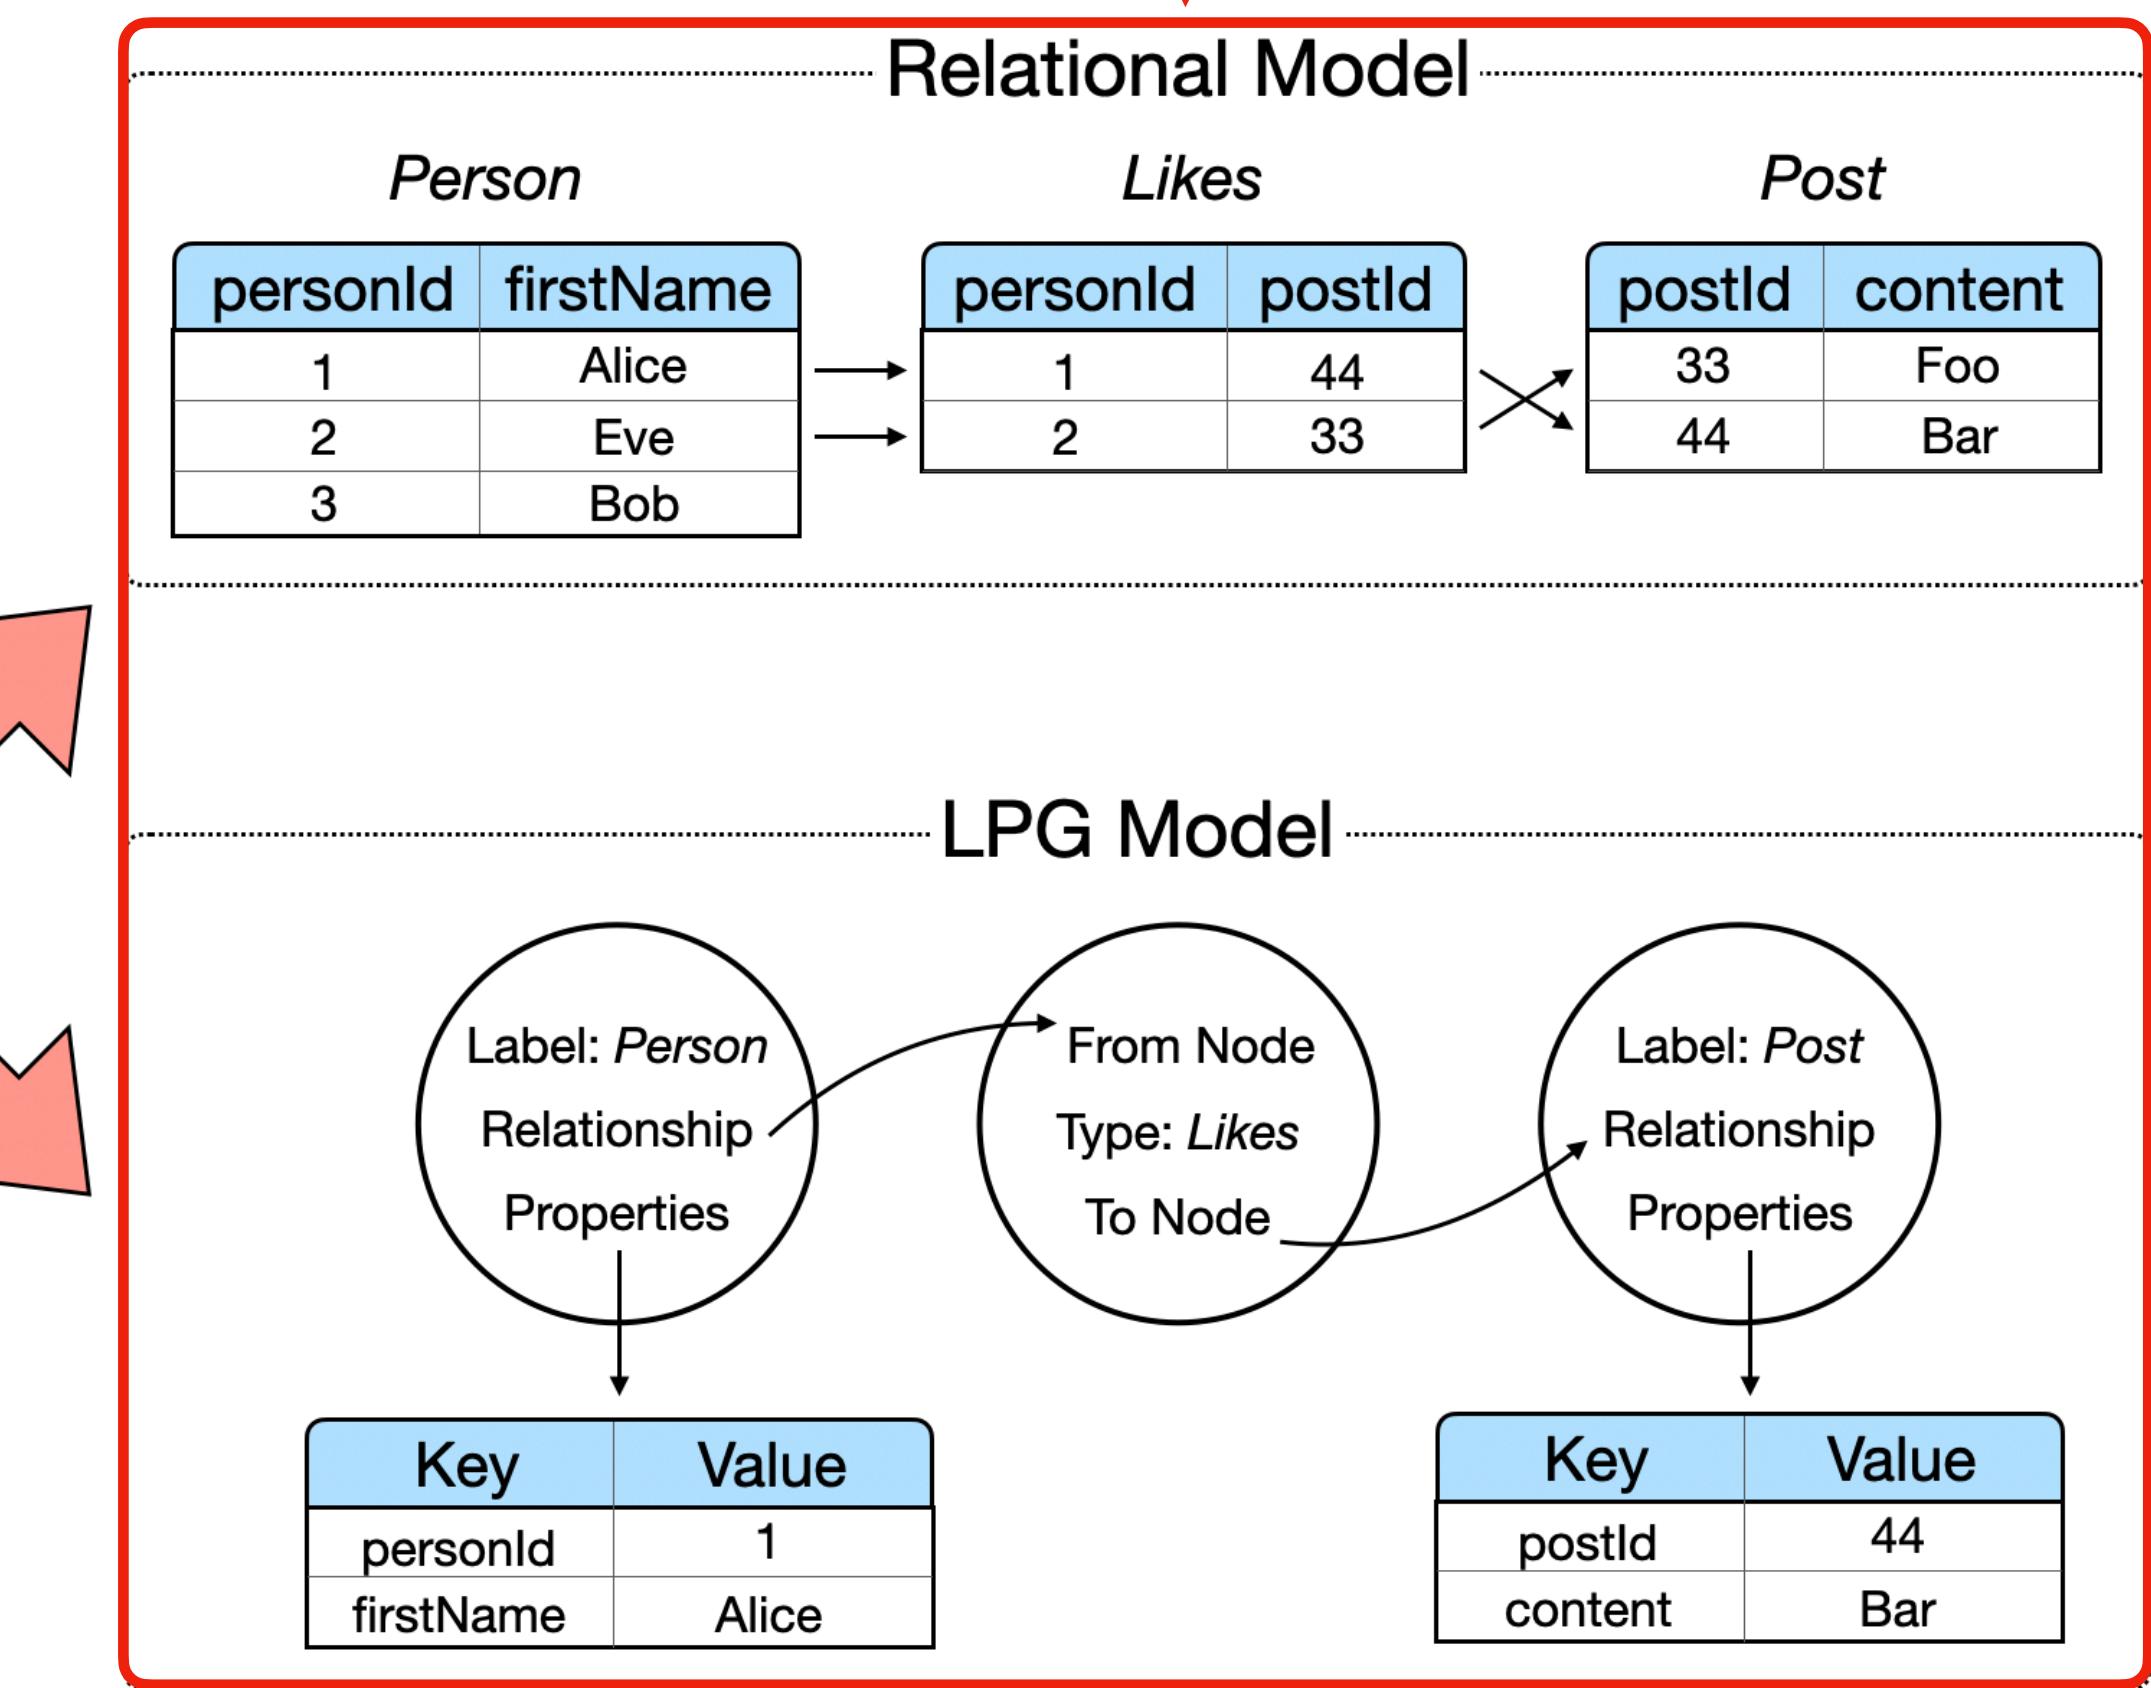
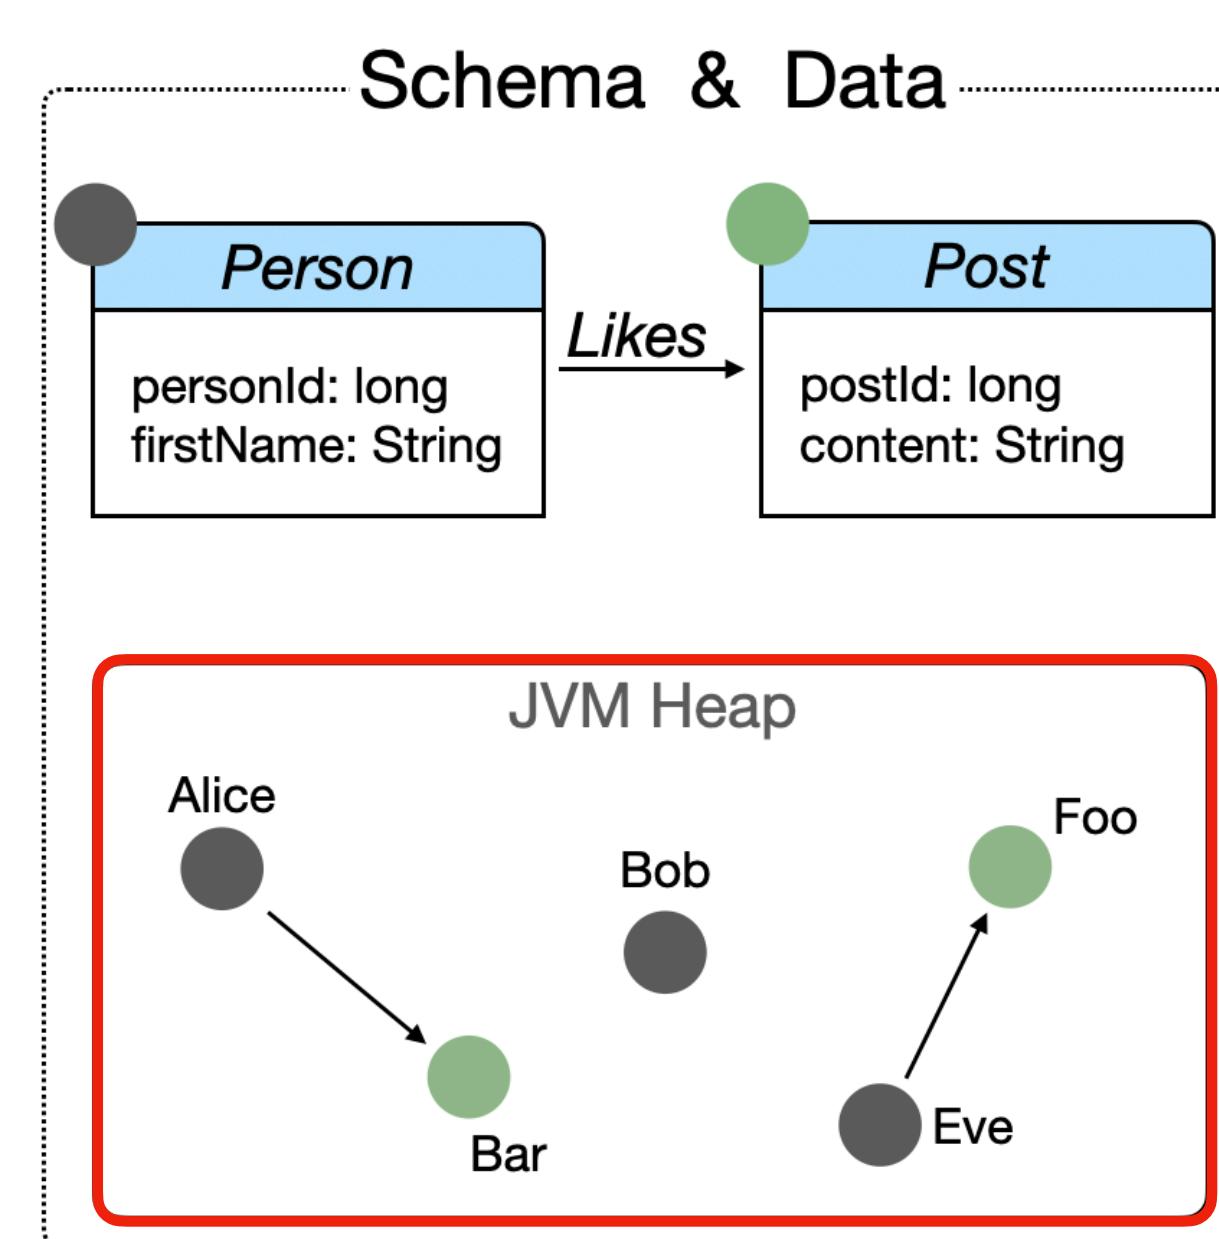
Motivation

Disconnect between storage models – data from objects has to be copied and stored in some form



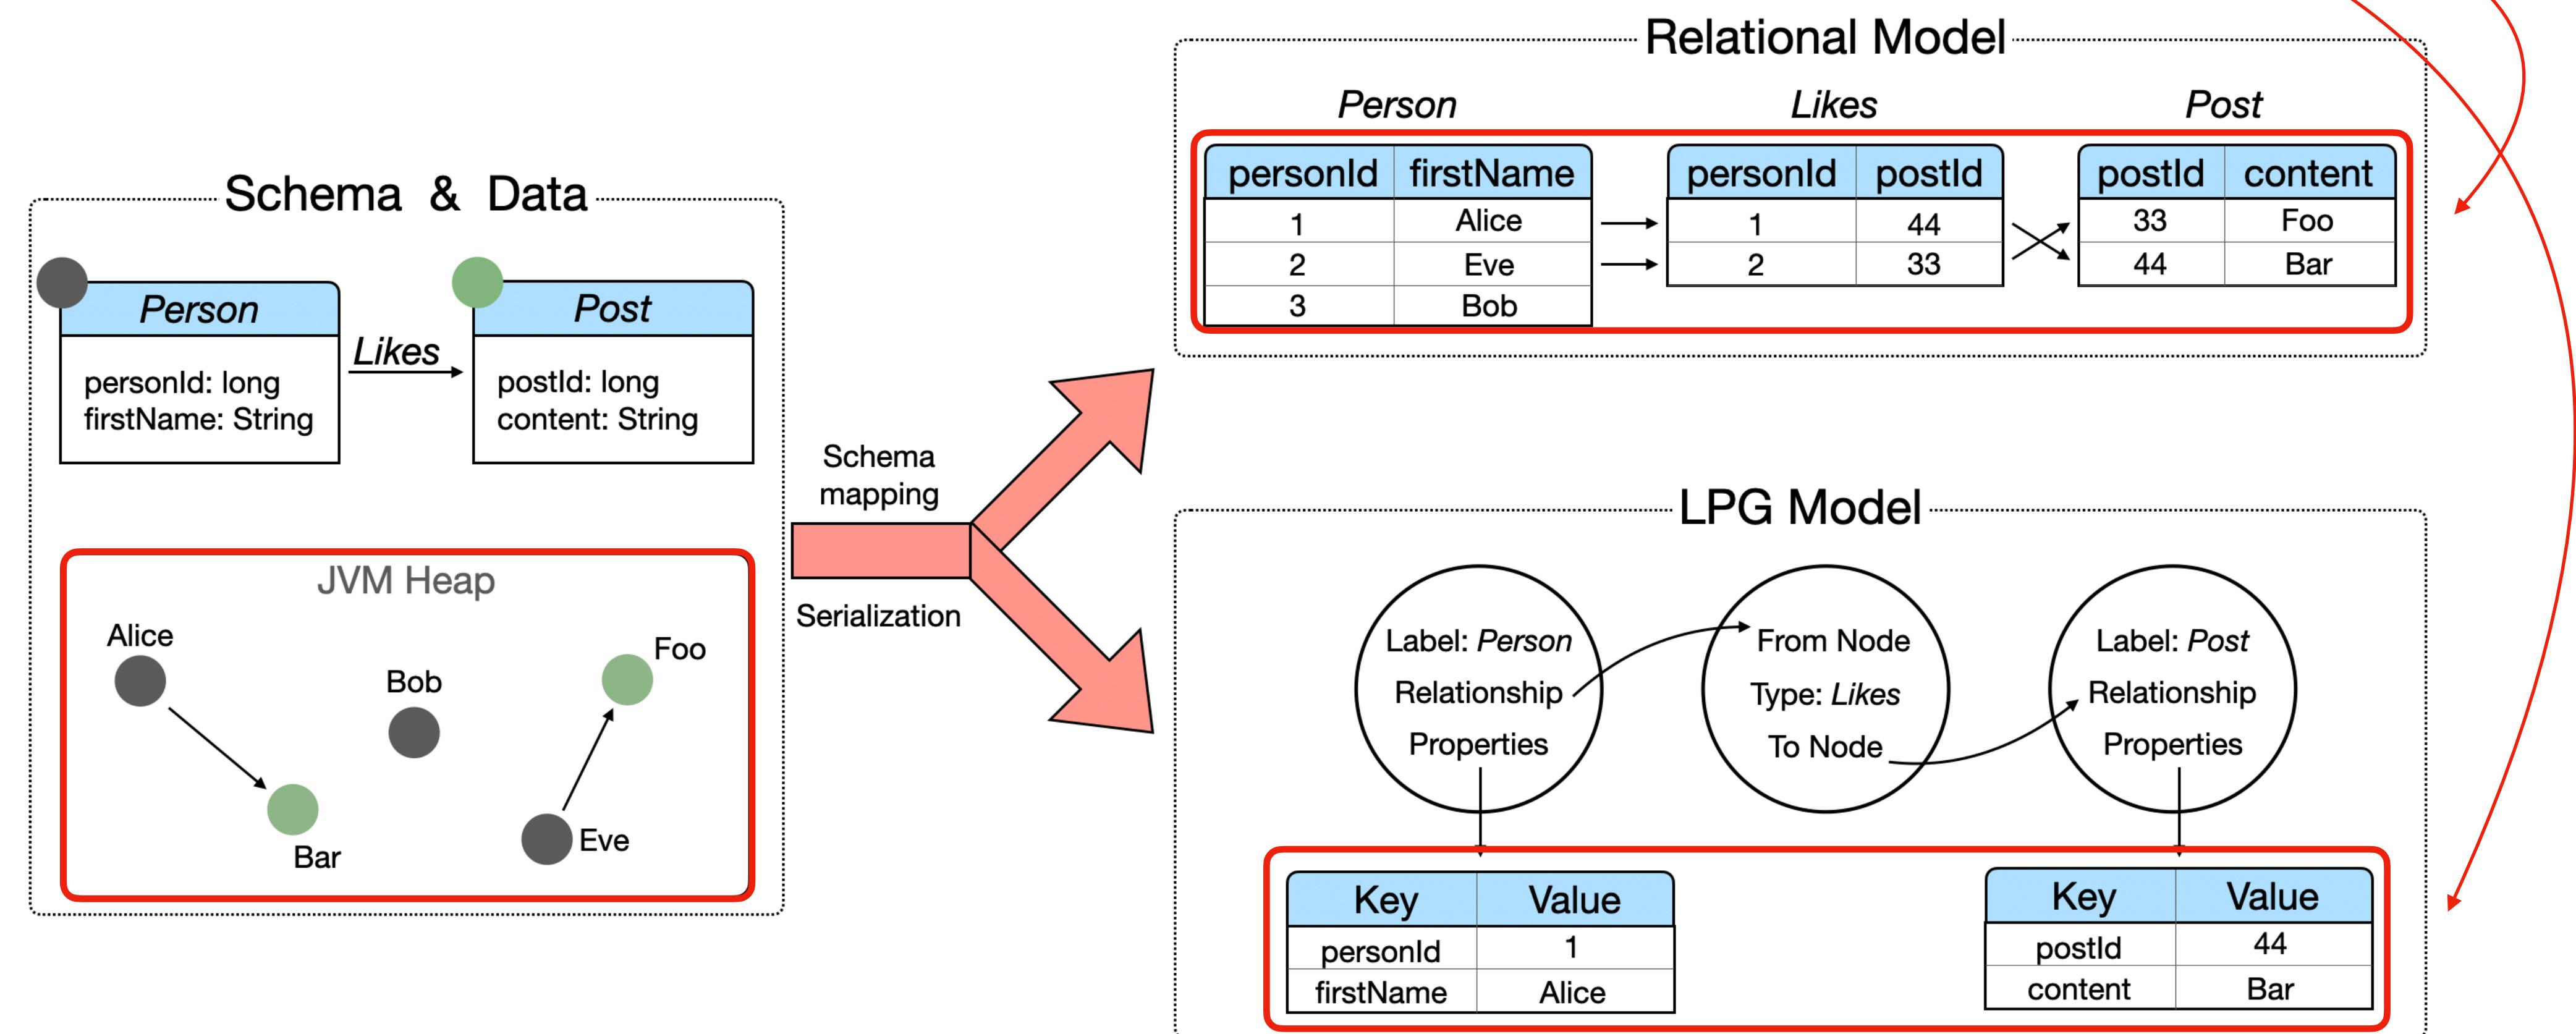
Motivation

Query edits have to be mapped backed to source
(not live)



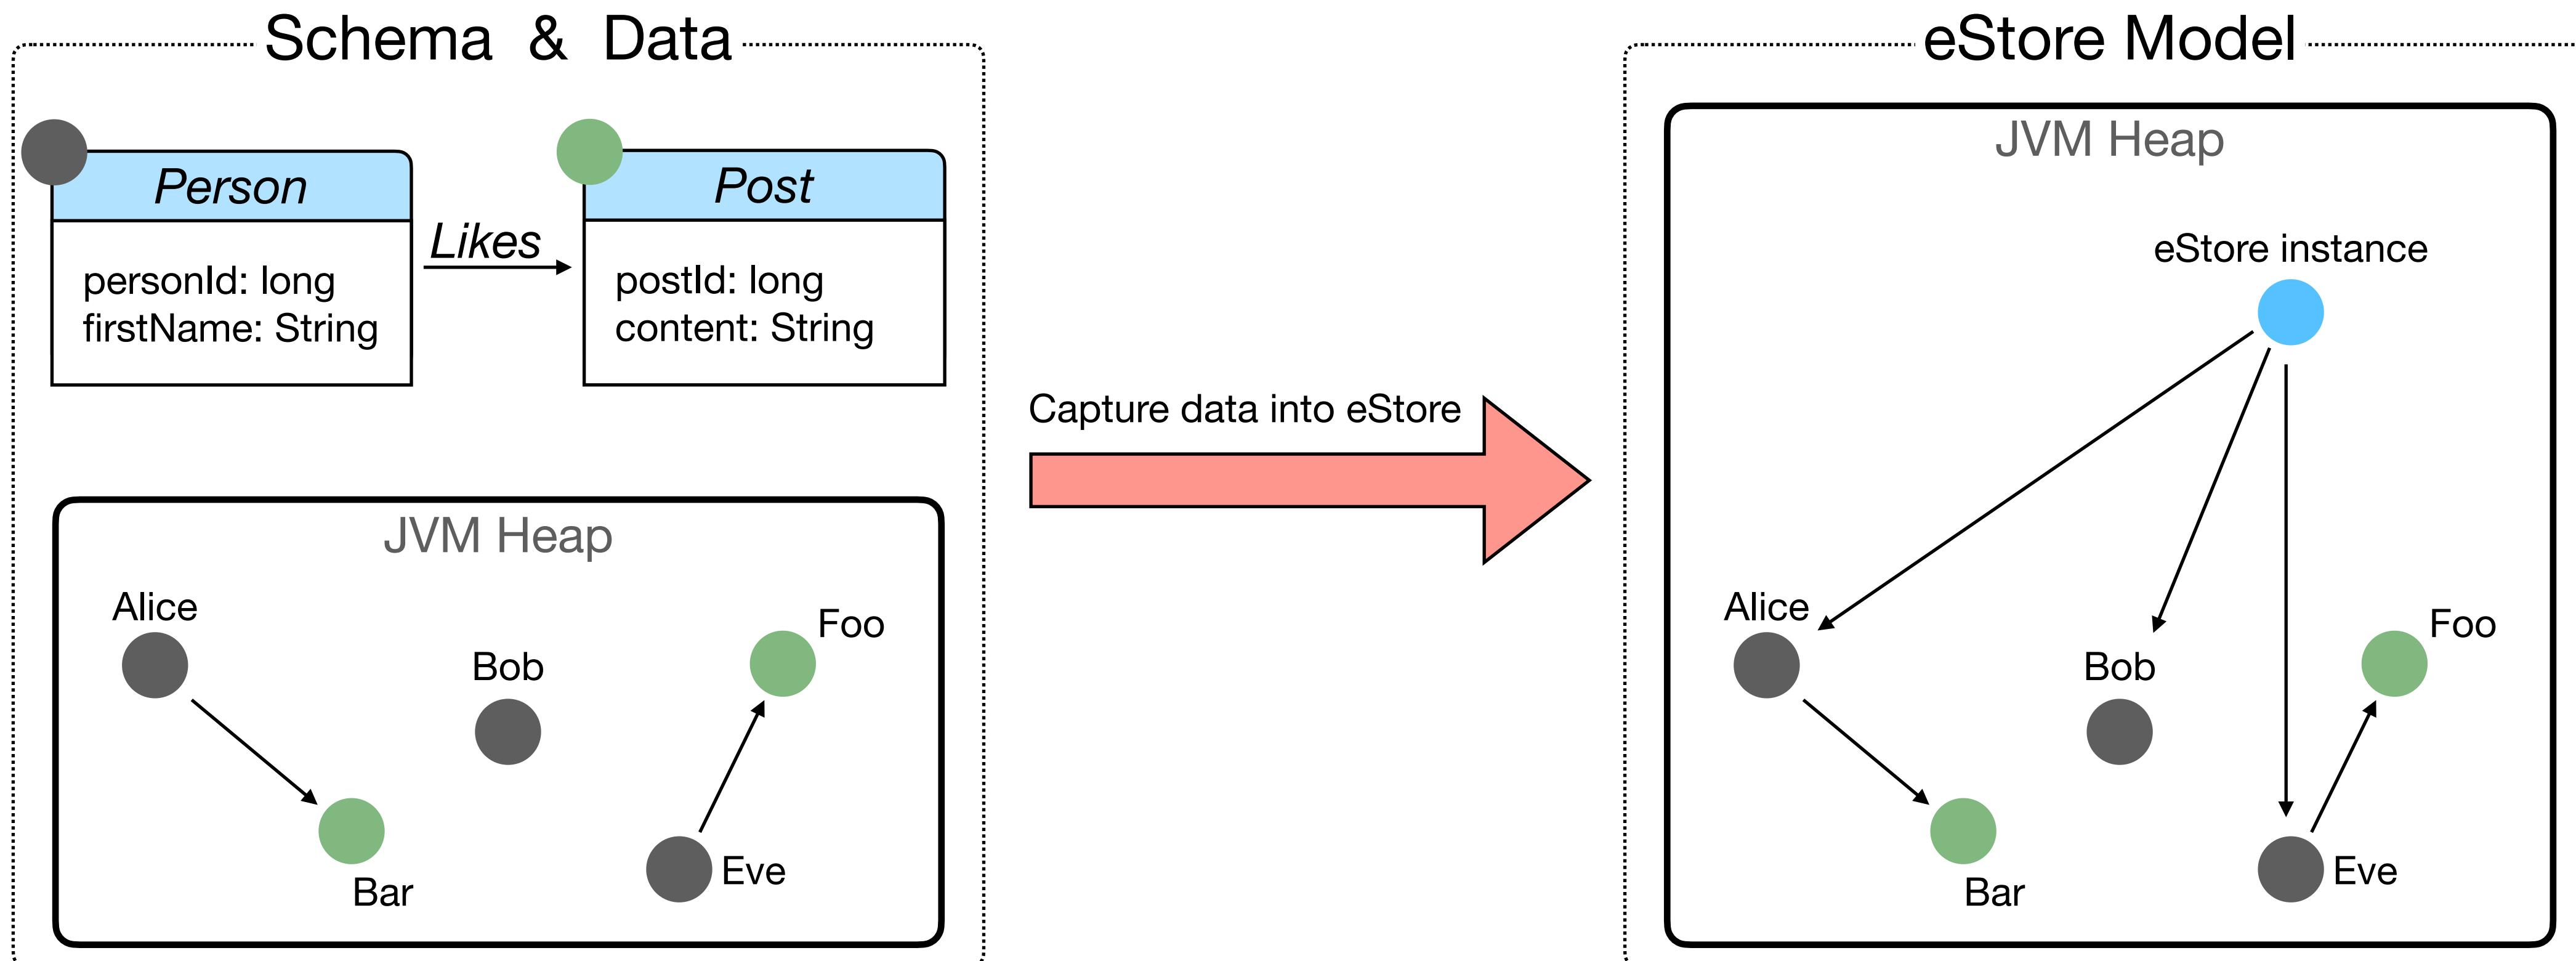
Motivation

Heavy abstraction and bloat – not ideal for quick setup and testing environments



Motivation

eStore is motivated by a simple observation: *your heap is already a richly connected graph—stop copying it, and start querying it directly.*



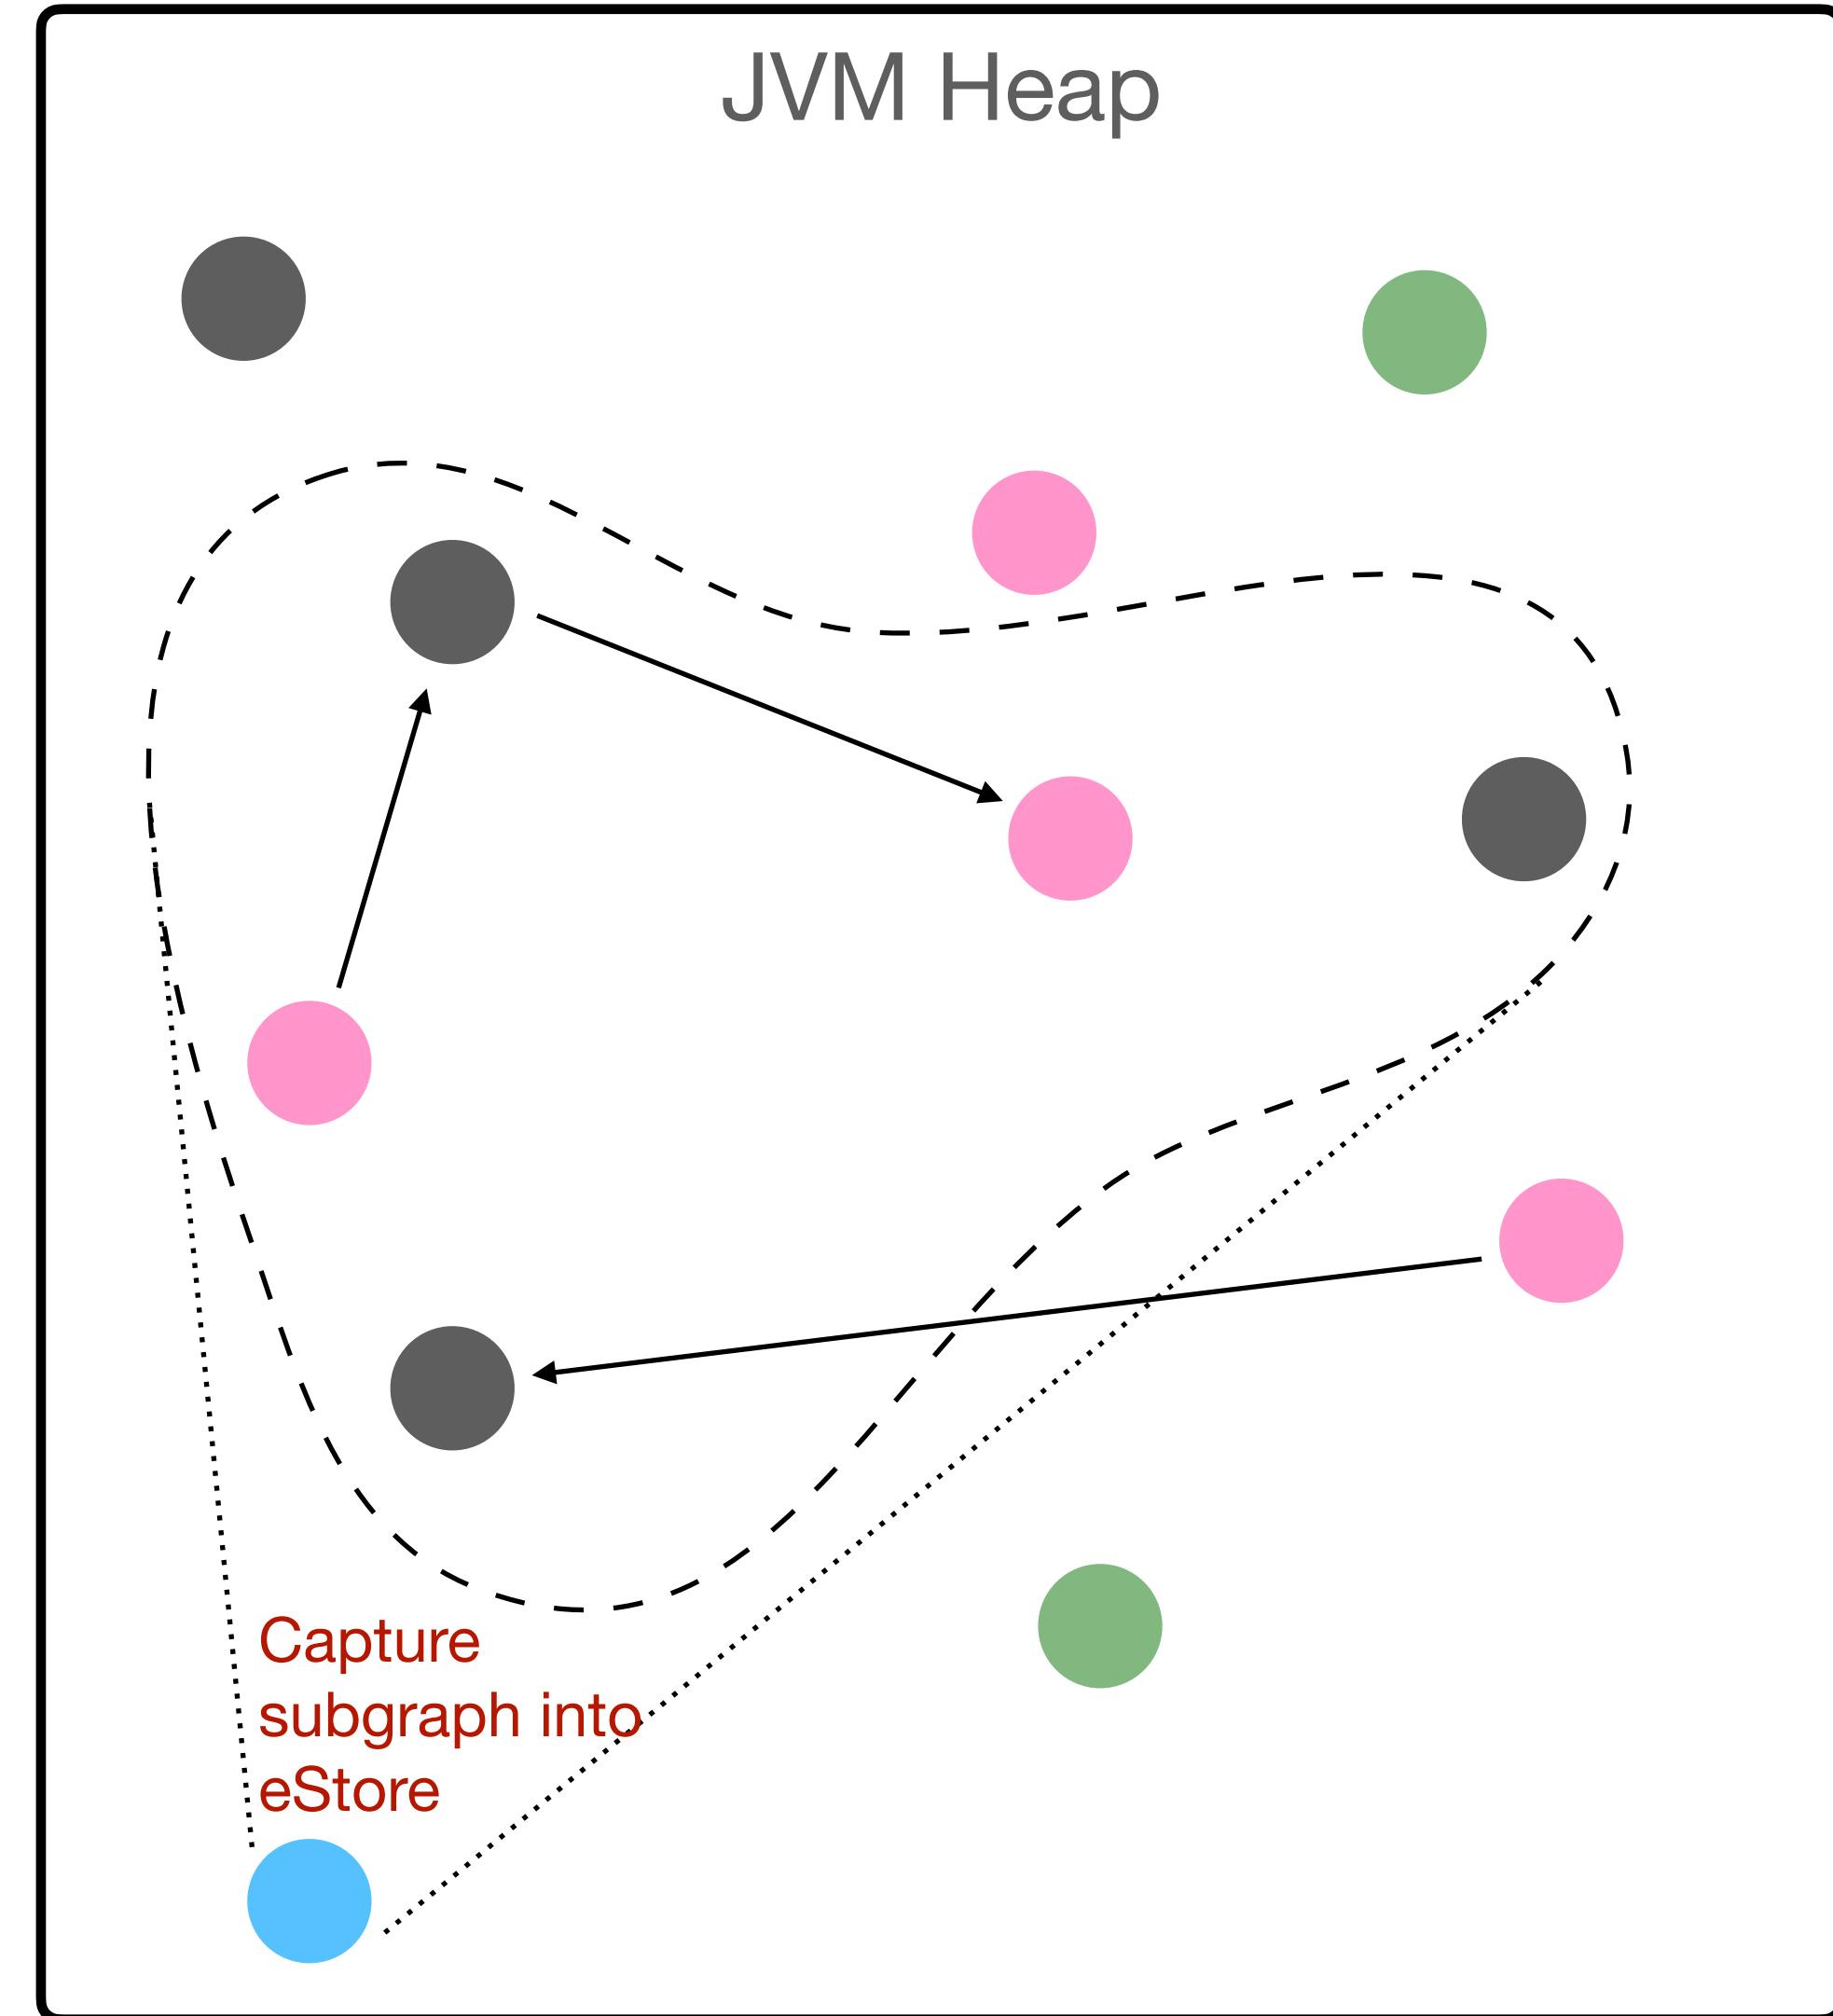
Our Contributions

- eStore: An in-memory object graph store that allows declarative queries over live objects and interoperability between imperative code and stored object
- Epsilon storage model: Allows storing objects with zero abstraction
- Compile-time query rewriting into imperative code to reduce runtime cost
- Formalization of the epsilon storage model
- Evaluation
 - Demonstrate robustness and lightweight design of eStore by evaluating on the LDBC SNB benchmark
 - Demonstrate eStore superior performance over existing heap querying frameworks (OGO)
 - Demonstrate the performance improvement with compile-time query rewriting

Insight behind eStore

- We see the heap of the JVM as a **graph** whose subgraphs can be captured into a graph object store and queried
 - Objects as nodes
 - References as edges
- We use expressive declarative graph query language **Cypher** to query the graph object store
- Example:

```
1| public void testAcyclicity(){  
2|     List<Long> list = new LinkedList<Long>();  
3|     ... //Add nodes to the list  
4|     EpsilonStore db = new EpsilonStore("db");  
5|     db.captureAll(list);  
6|     assertTrue(db.query(  
7|         "MATCH (n:'LinkedList$Node')-[:next*]->(n)"  
8|         +" RETURN COUNT(n) = 0").getBoolean(0));  
9|}
```



Example | Creating eStore instance



L
E
G
E
N
D

	Class instance
	ArrayList instance
	HashMap instance
	eStore instance
	Post instance
	Person instance
	instanceof relation
	Key
	List element relation

```
1| EpsilonStore db = new EpsilonStore ("dbname");
```

Example | Creating eStore instance

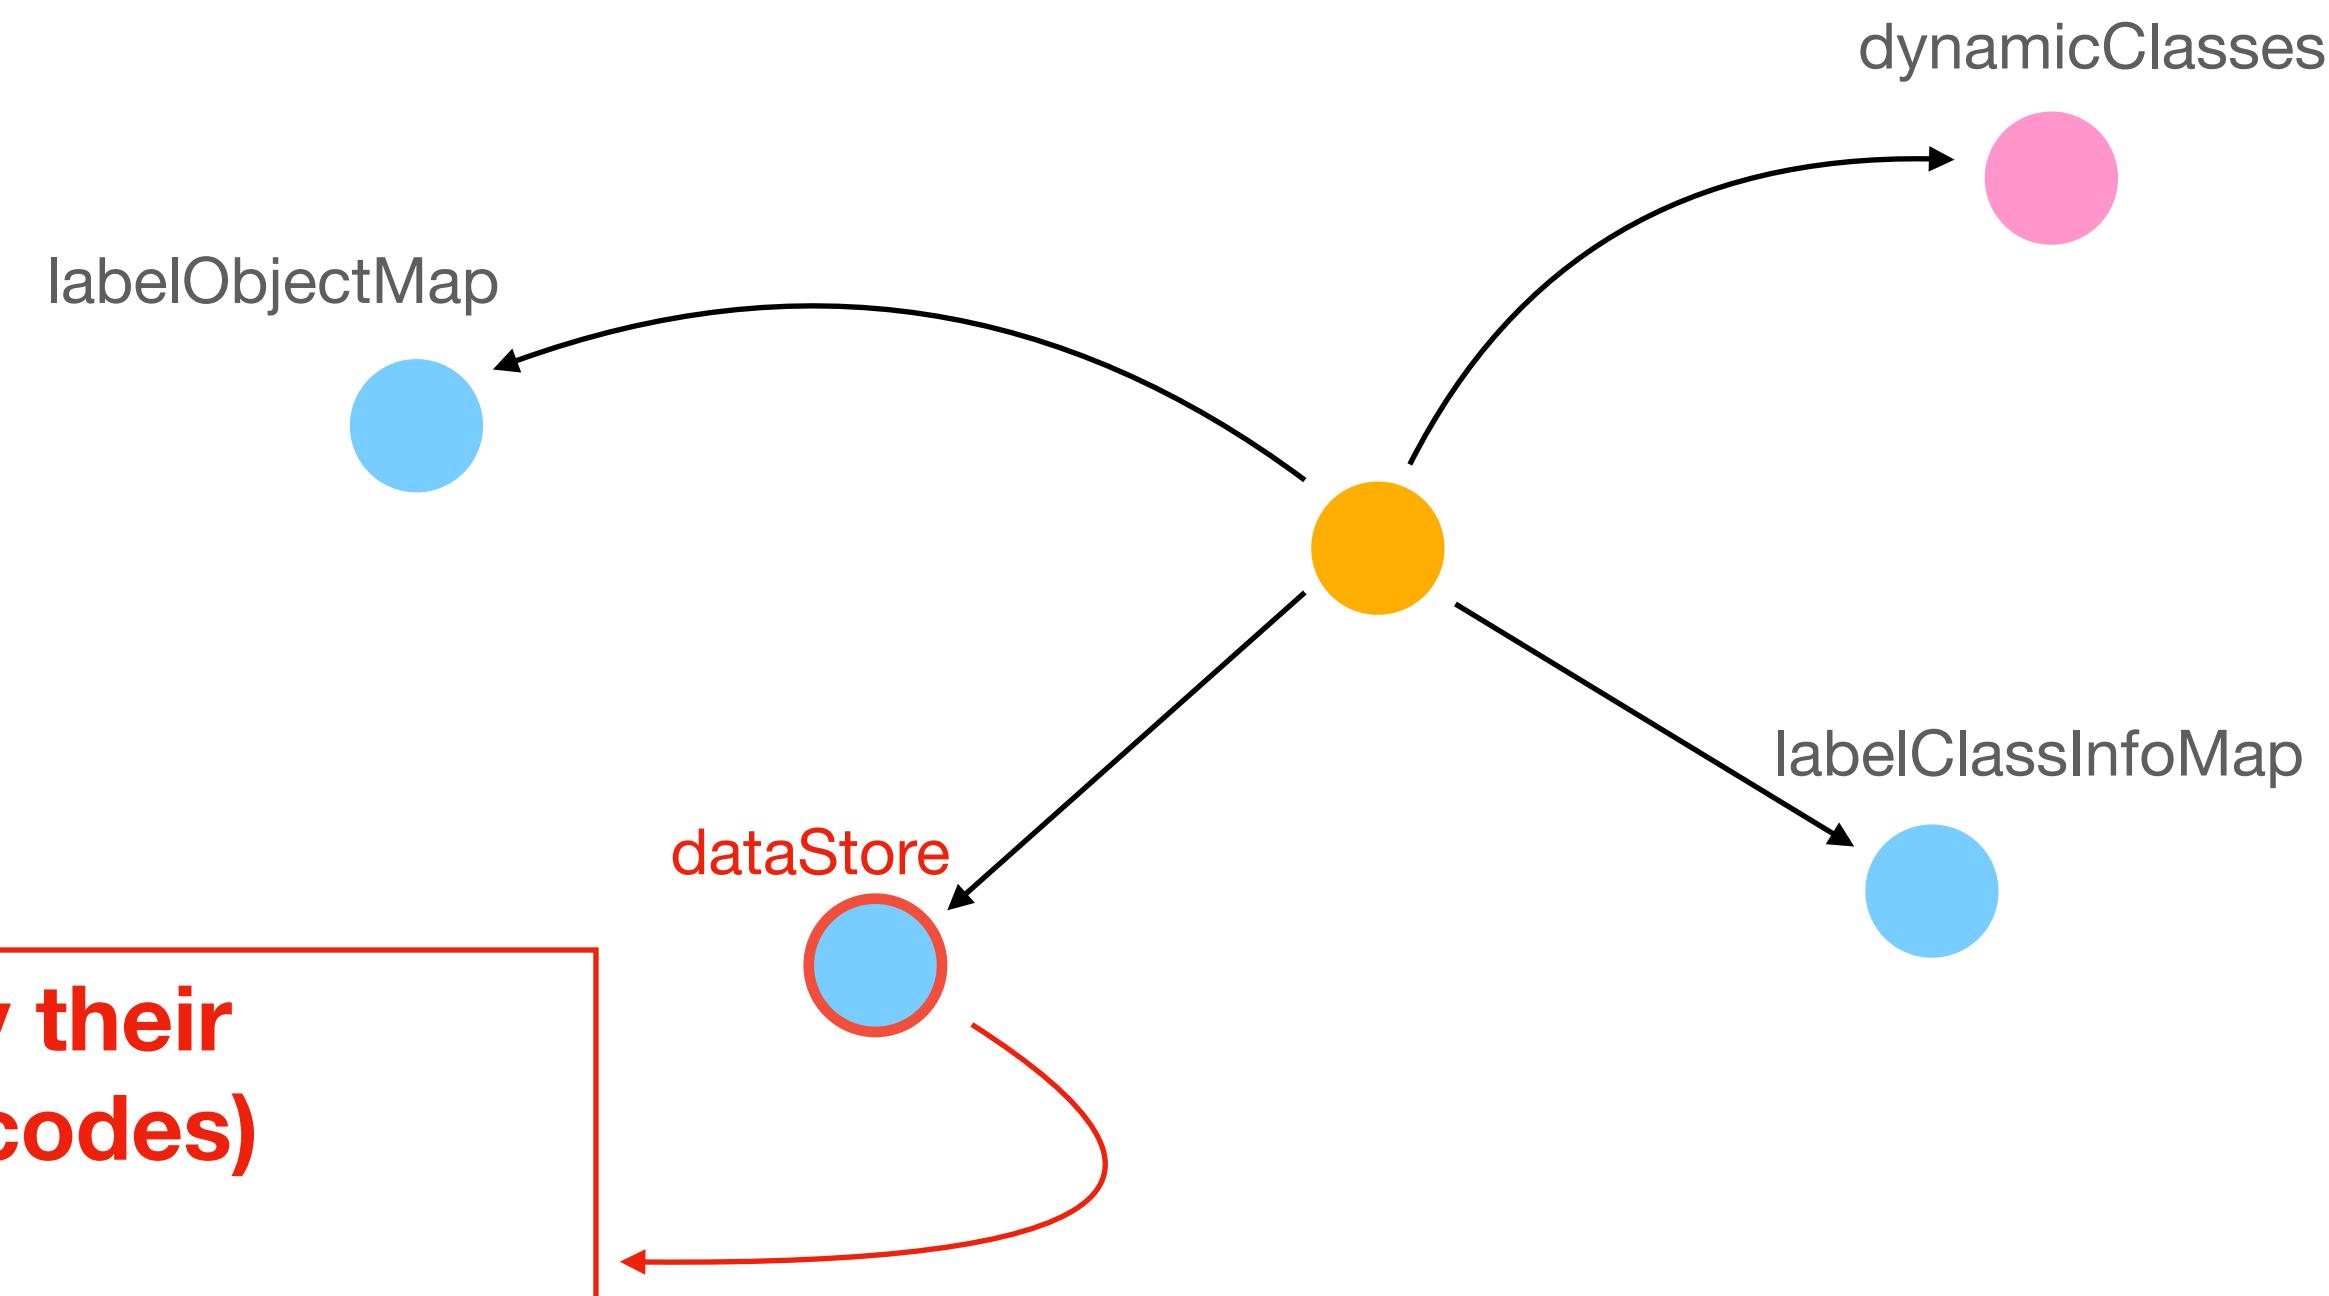


Stores objects by their class names (labels)
=>
Optimizes queries by restricting scans to stored objects belonging to label

```
1| EpsilonStore db = new EpsilonStore ("dbname");
```



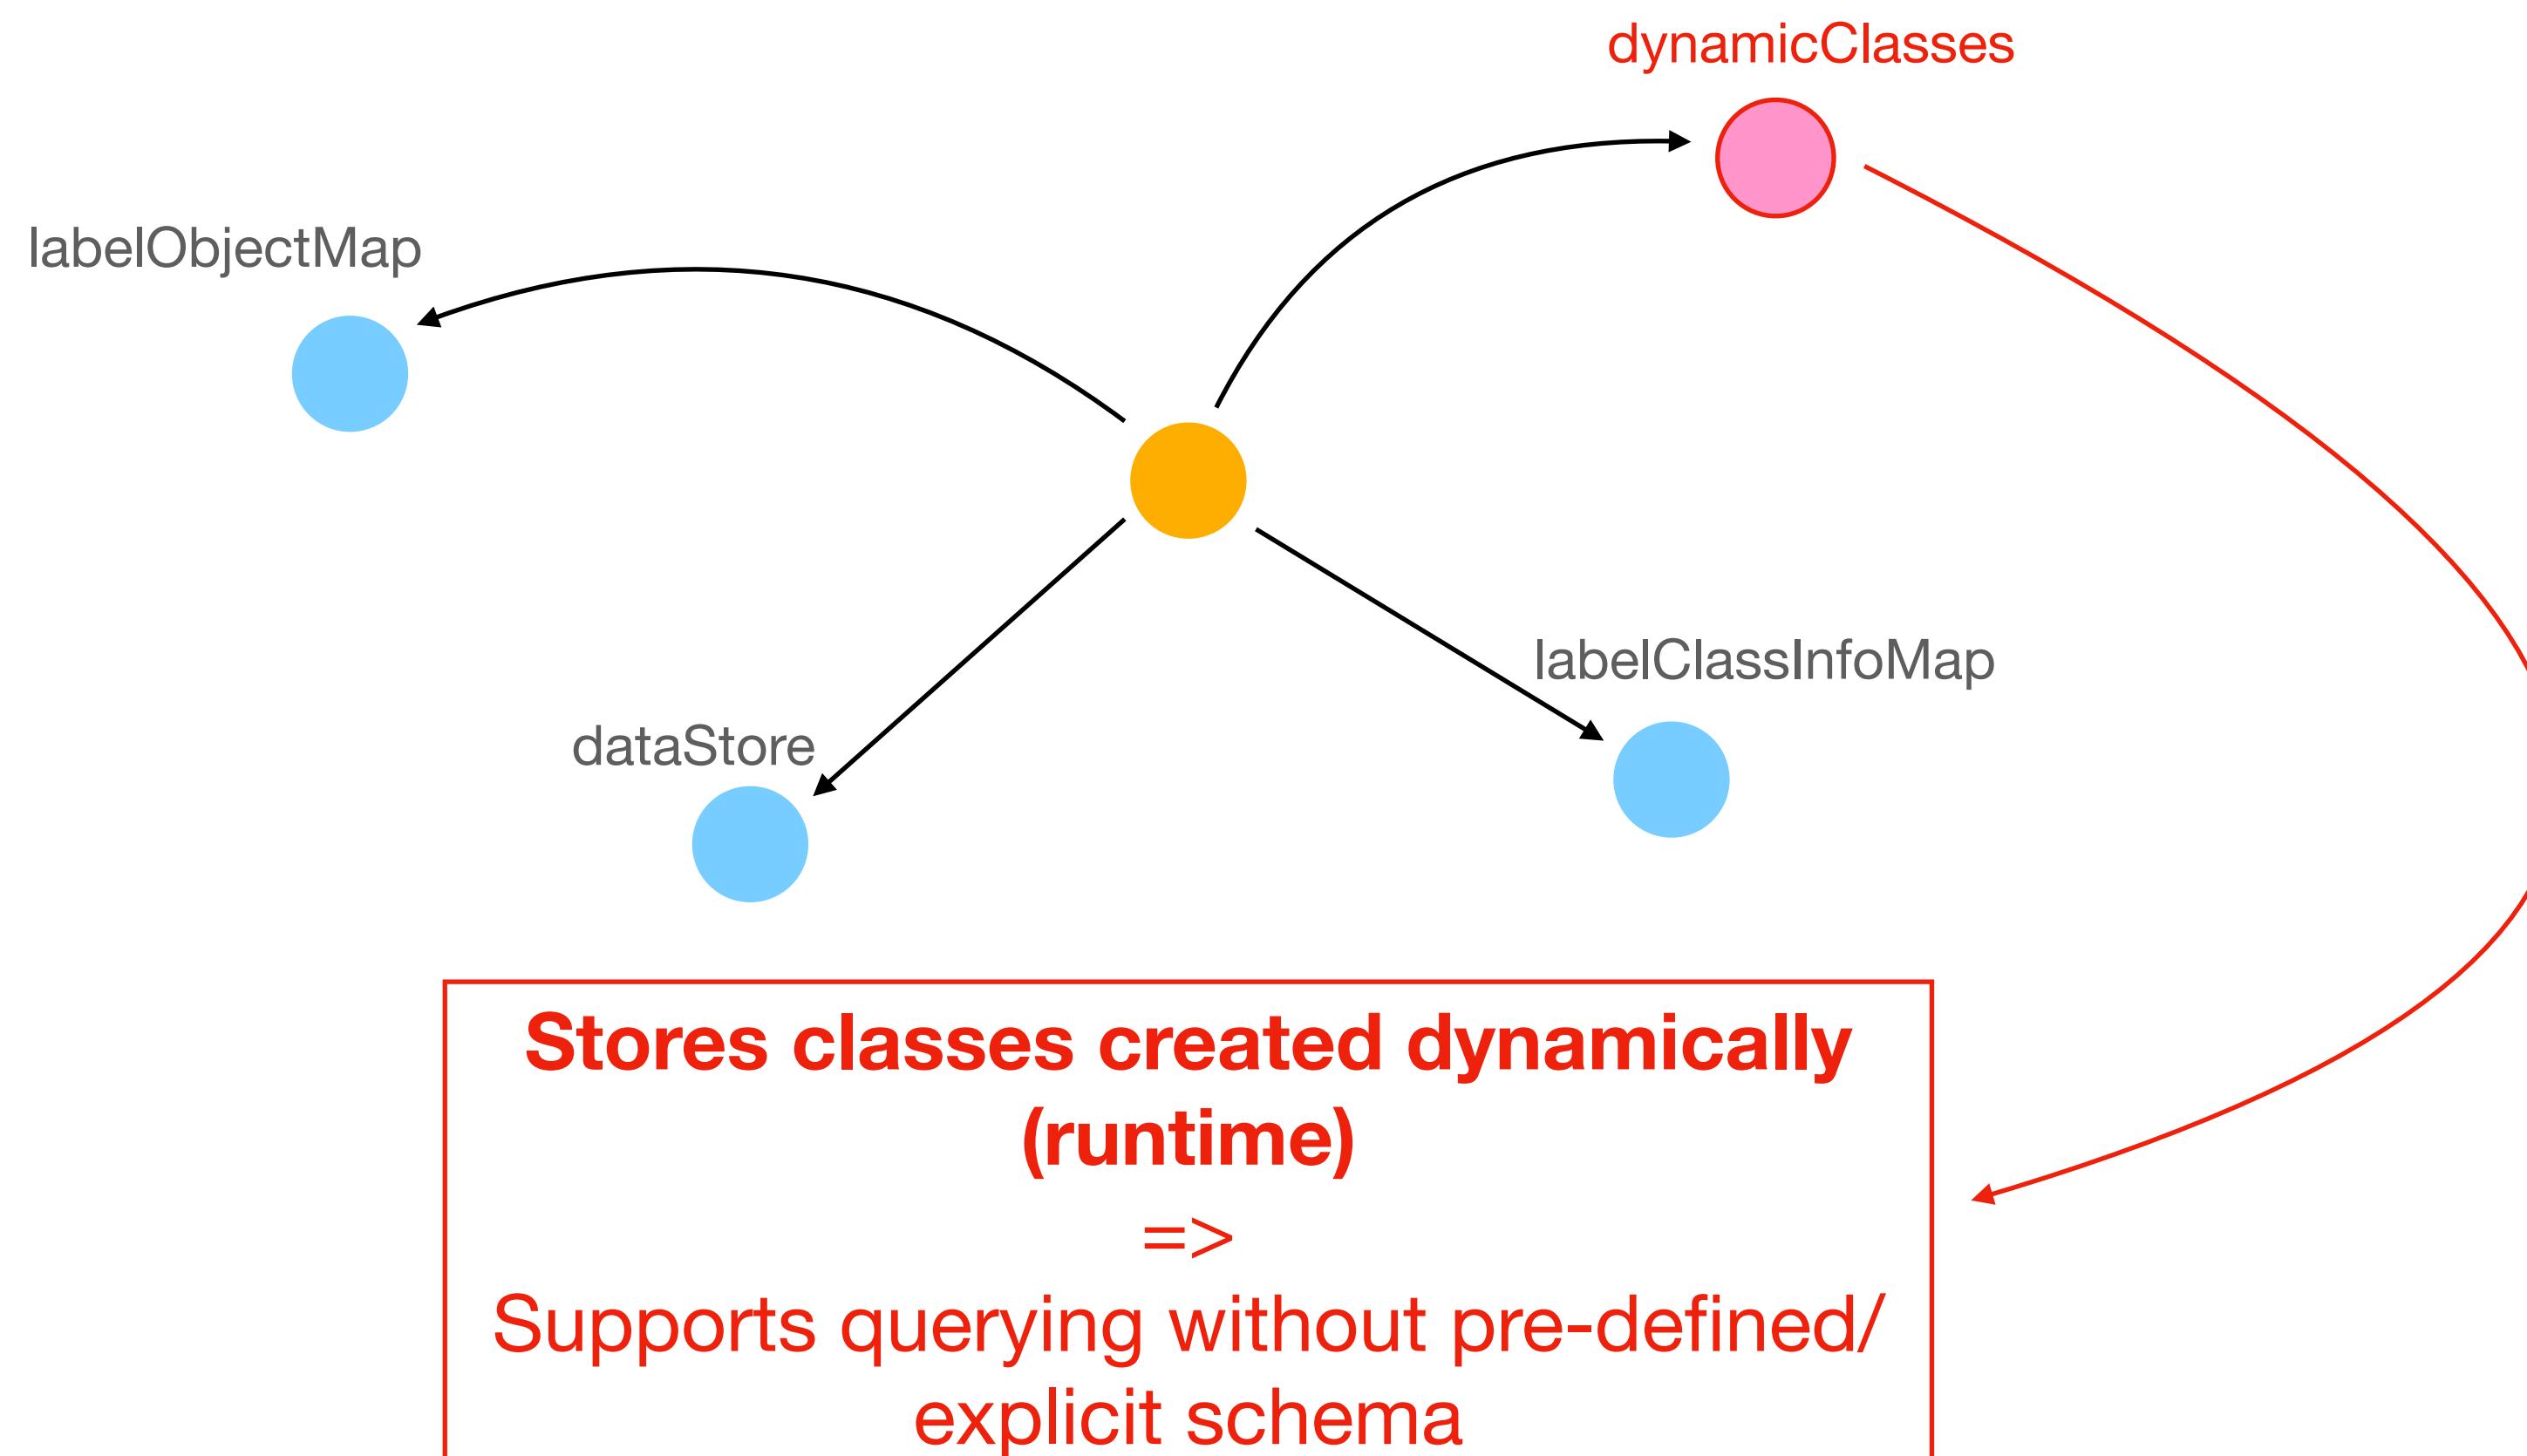
Example | Creating eStore instance



```
1| EpsilonStore db = new EpsilonStore ("dbname");
```



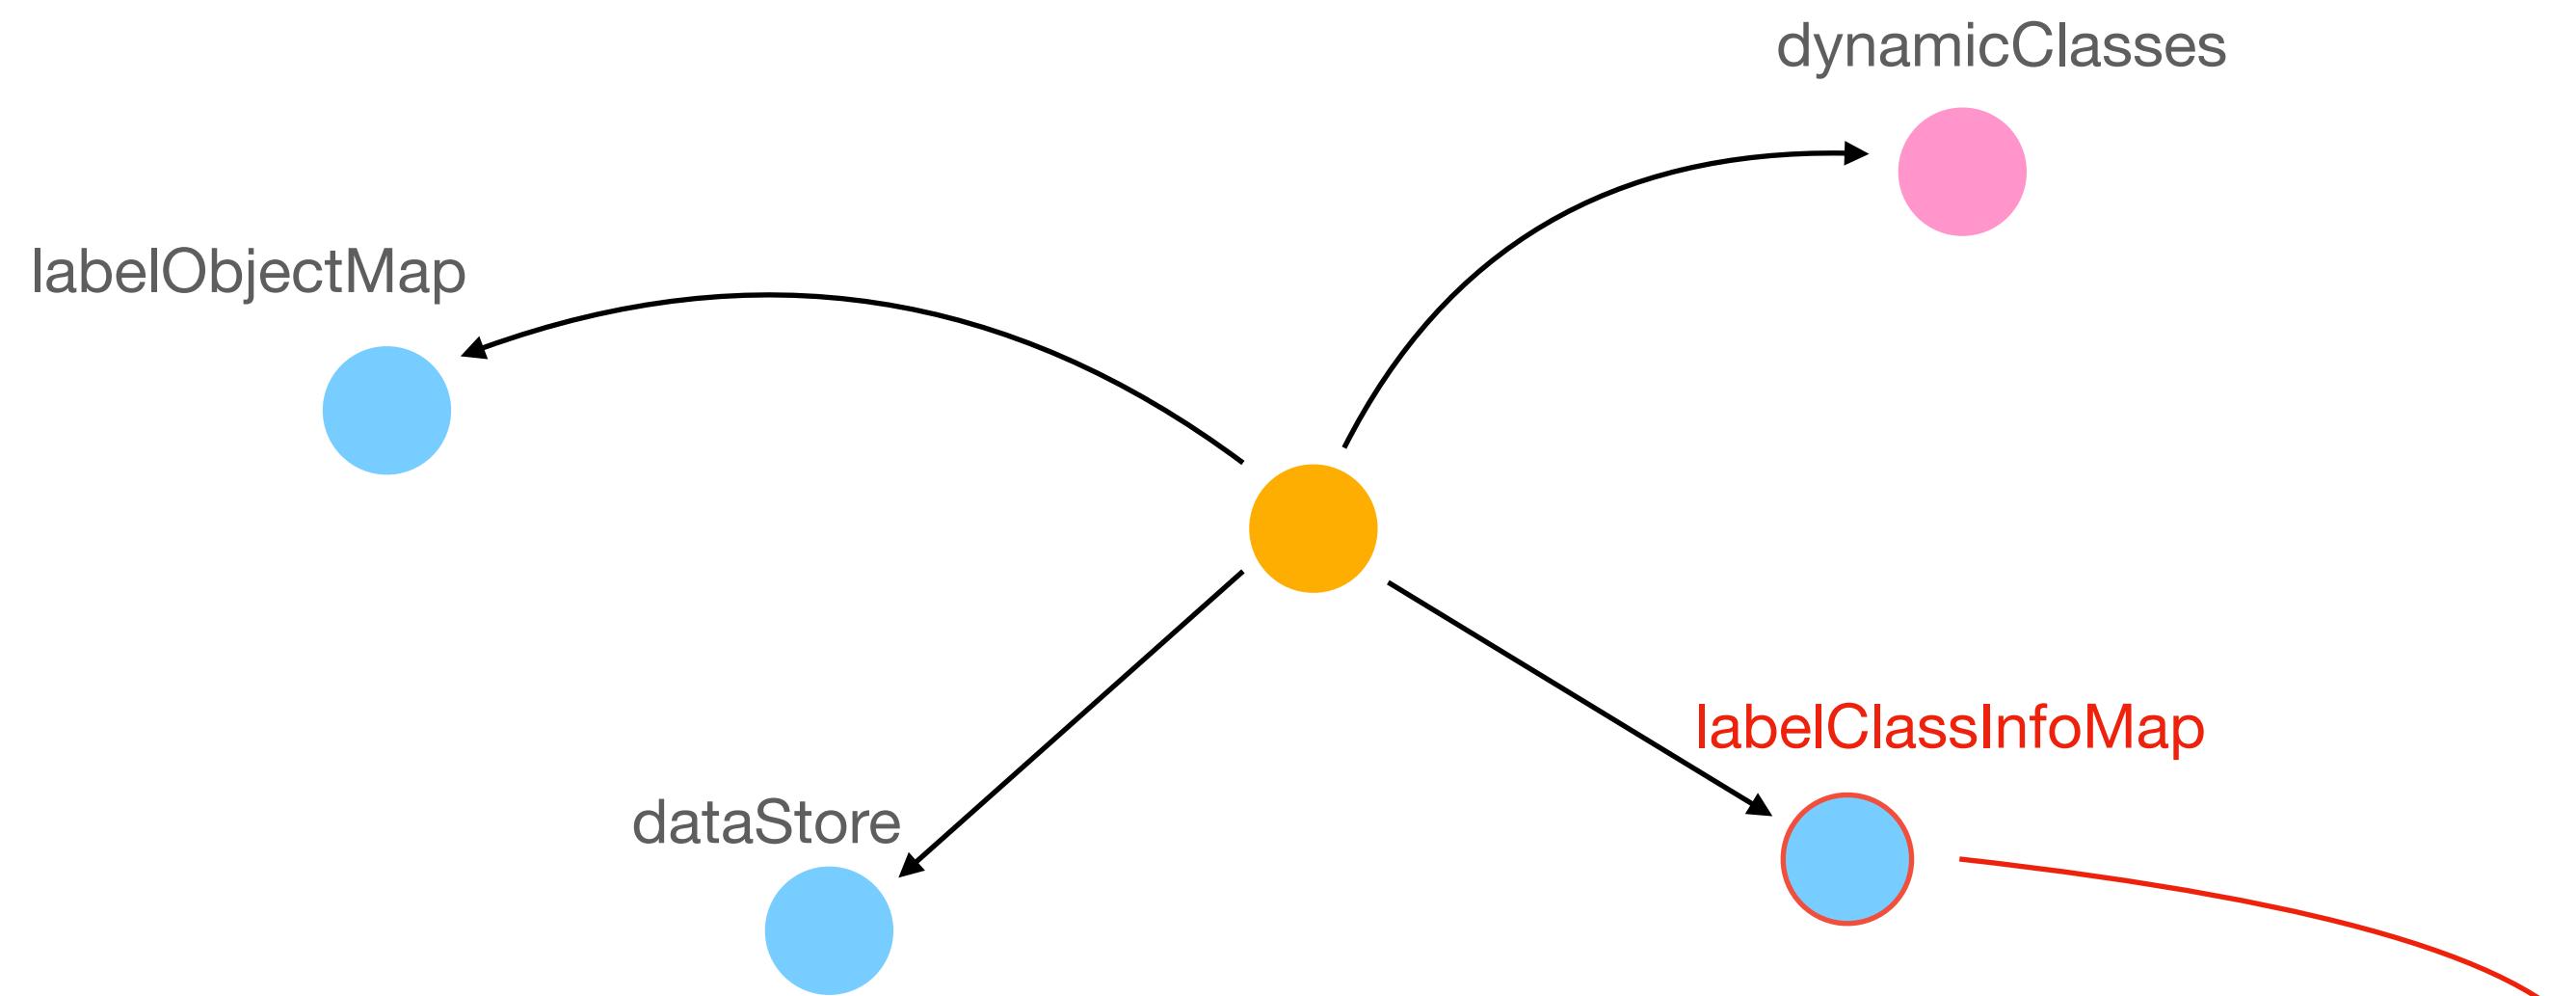
Example | Creating eStore instance



```
1| EpsilonStore db = new EpsilonStore ("dbname");
```



Example | Creating eStore instance



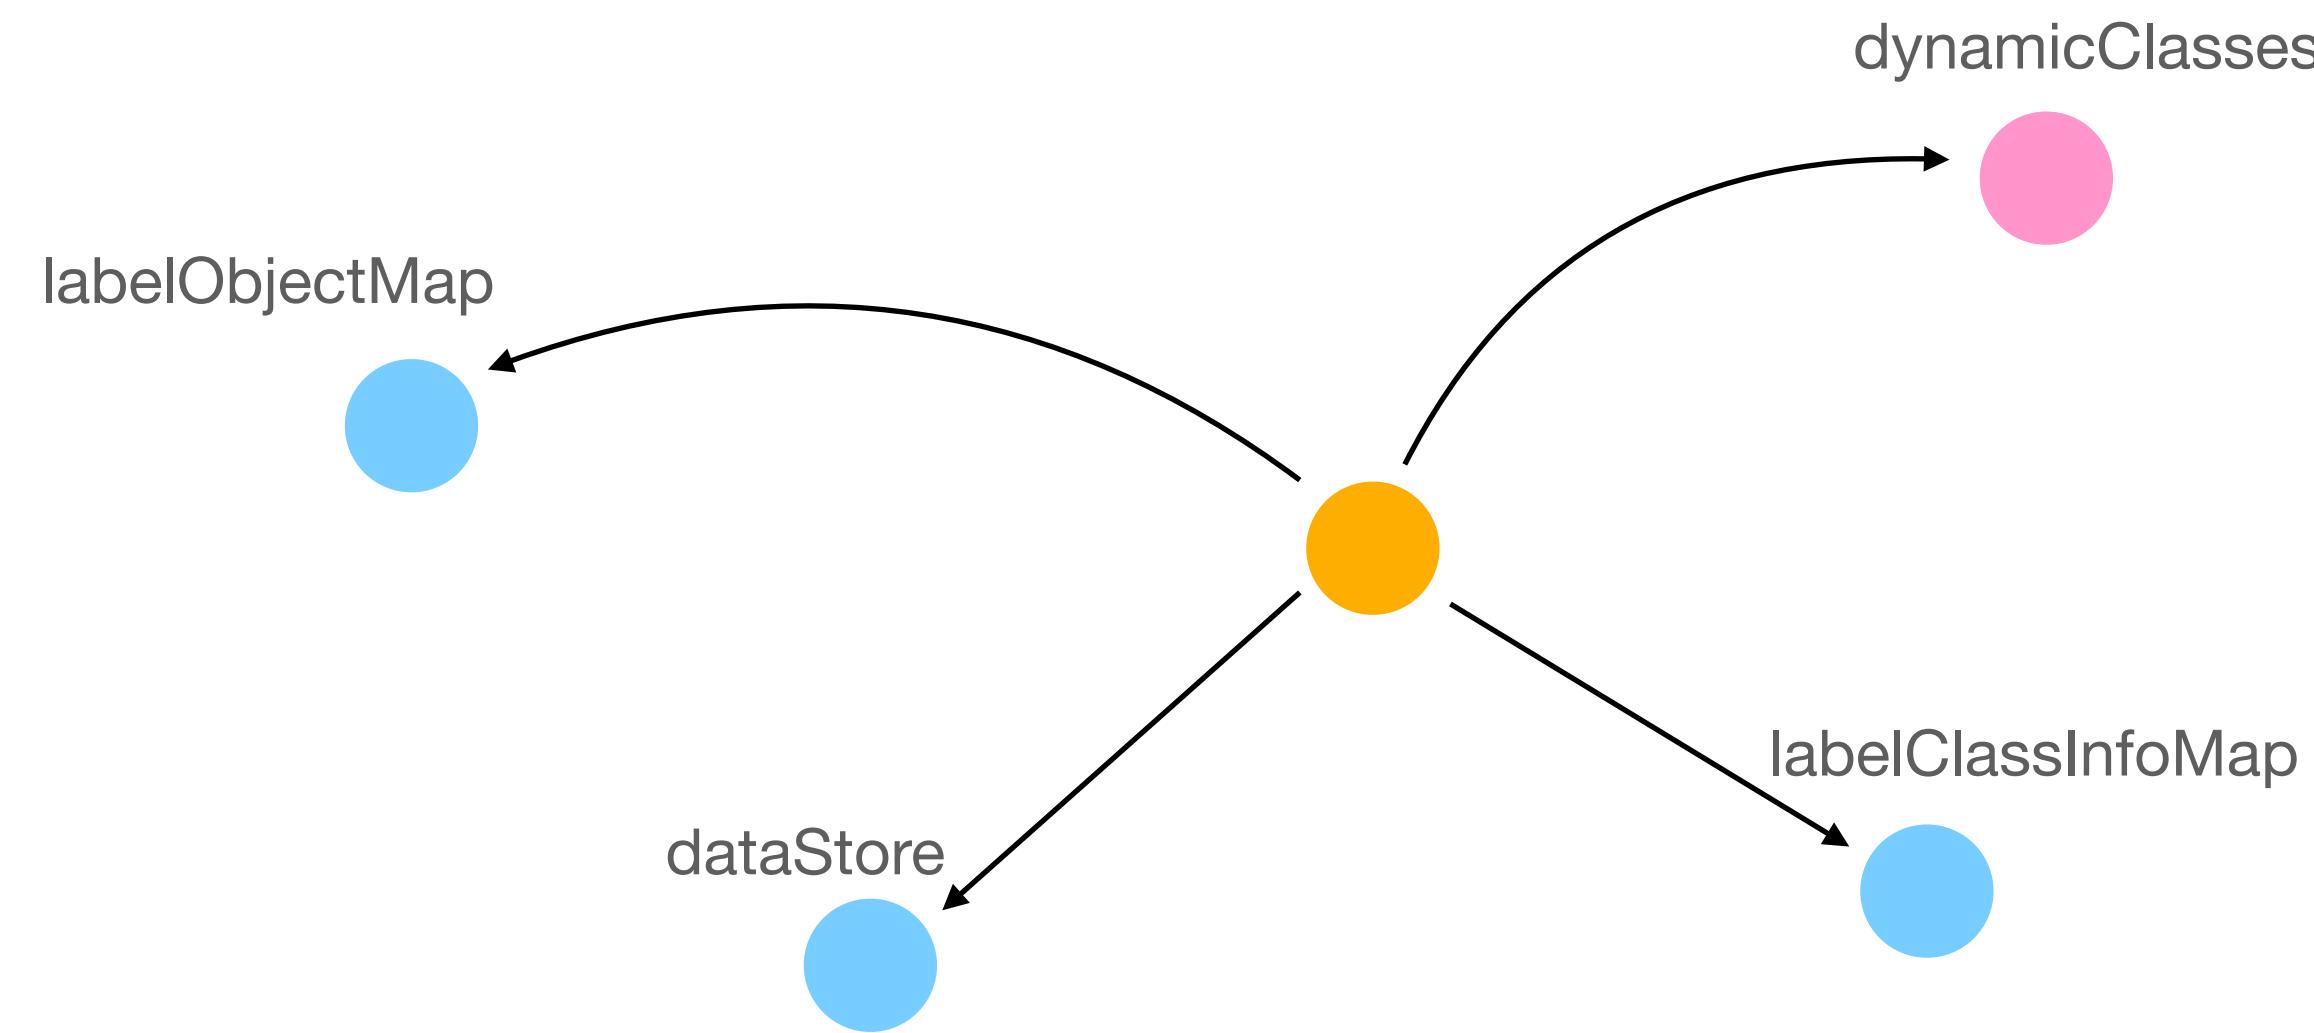
Maps classes of stored objects to their meta prop objects (ClassInfo)

=>
Optimizing queries containing node properties by caching class field names and types

```
1| EpsilonStore db = new EpsilonStore ("dbname");
```



Example | Adding data without pre-defined schema

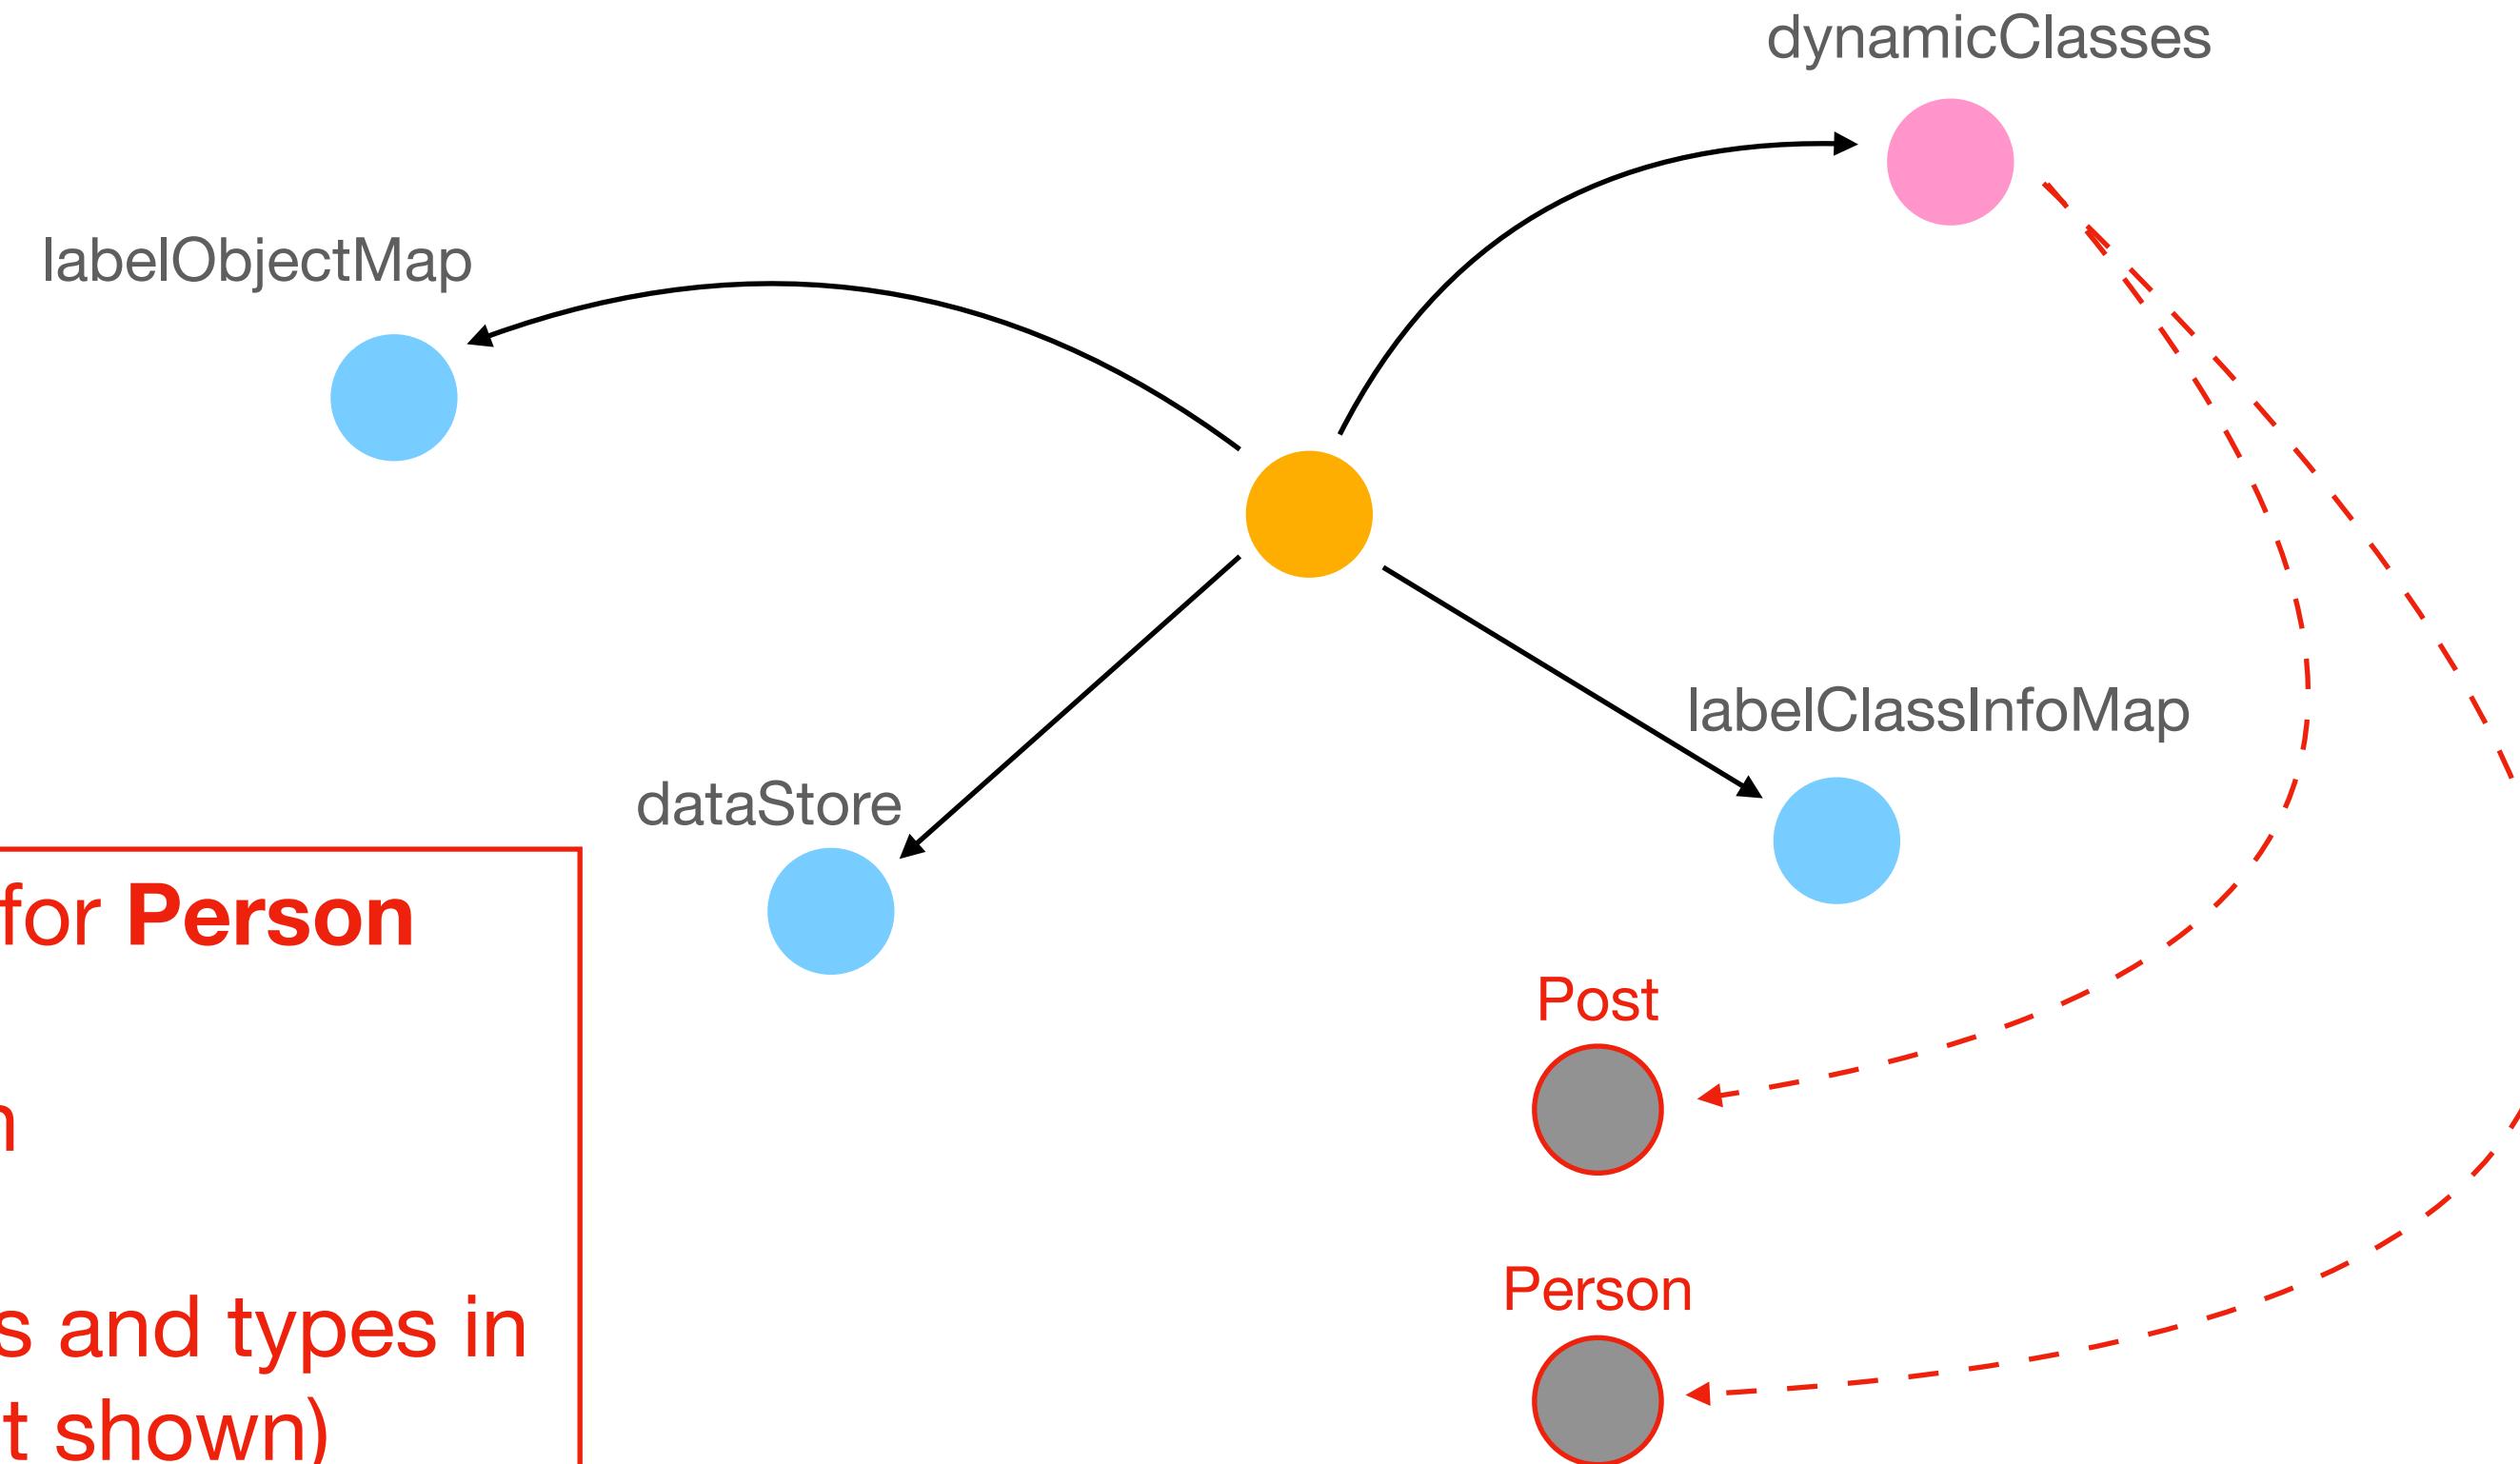


```
1| db.query(  
2| "CREATE (m:`Person` {id: 112, firstName: 'Eve'})"  
3| +"-[r:LIKES]->"  
4| +"(n:`Post` {id: 481, content: 'About Databases'})"  
5|);
```



Example | Adding data without pre-defined schema

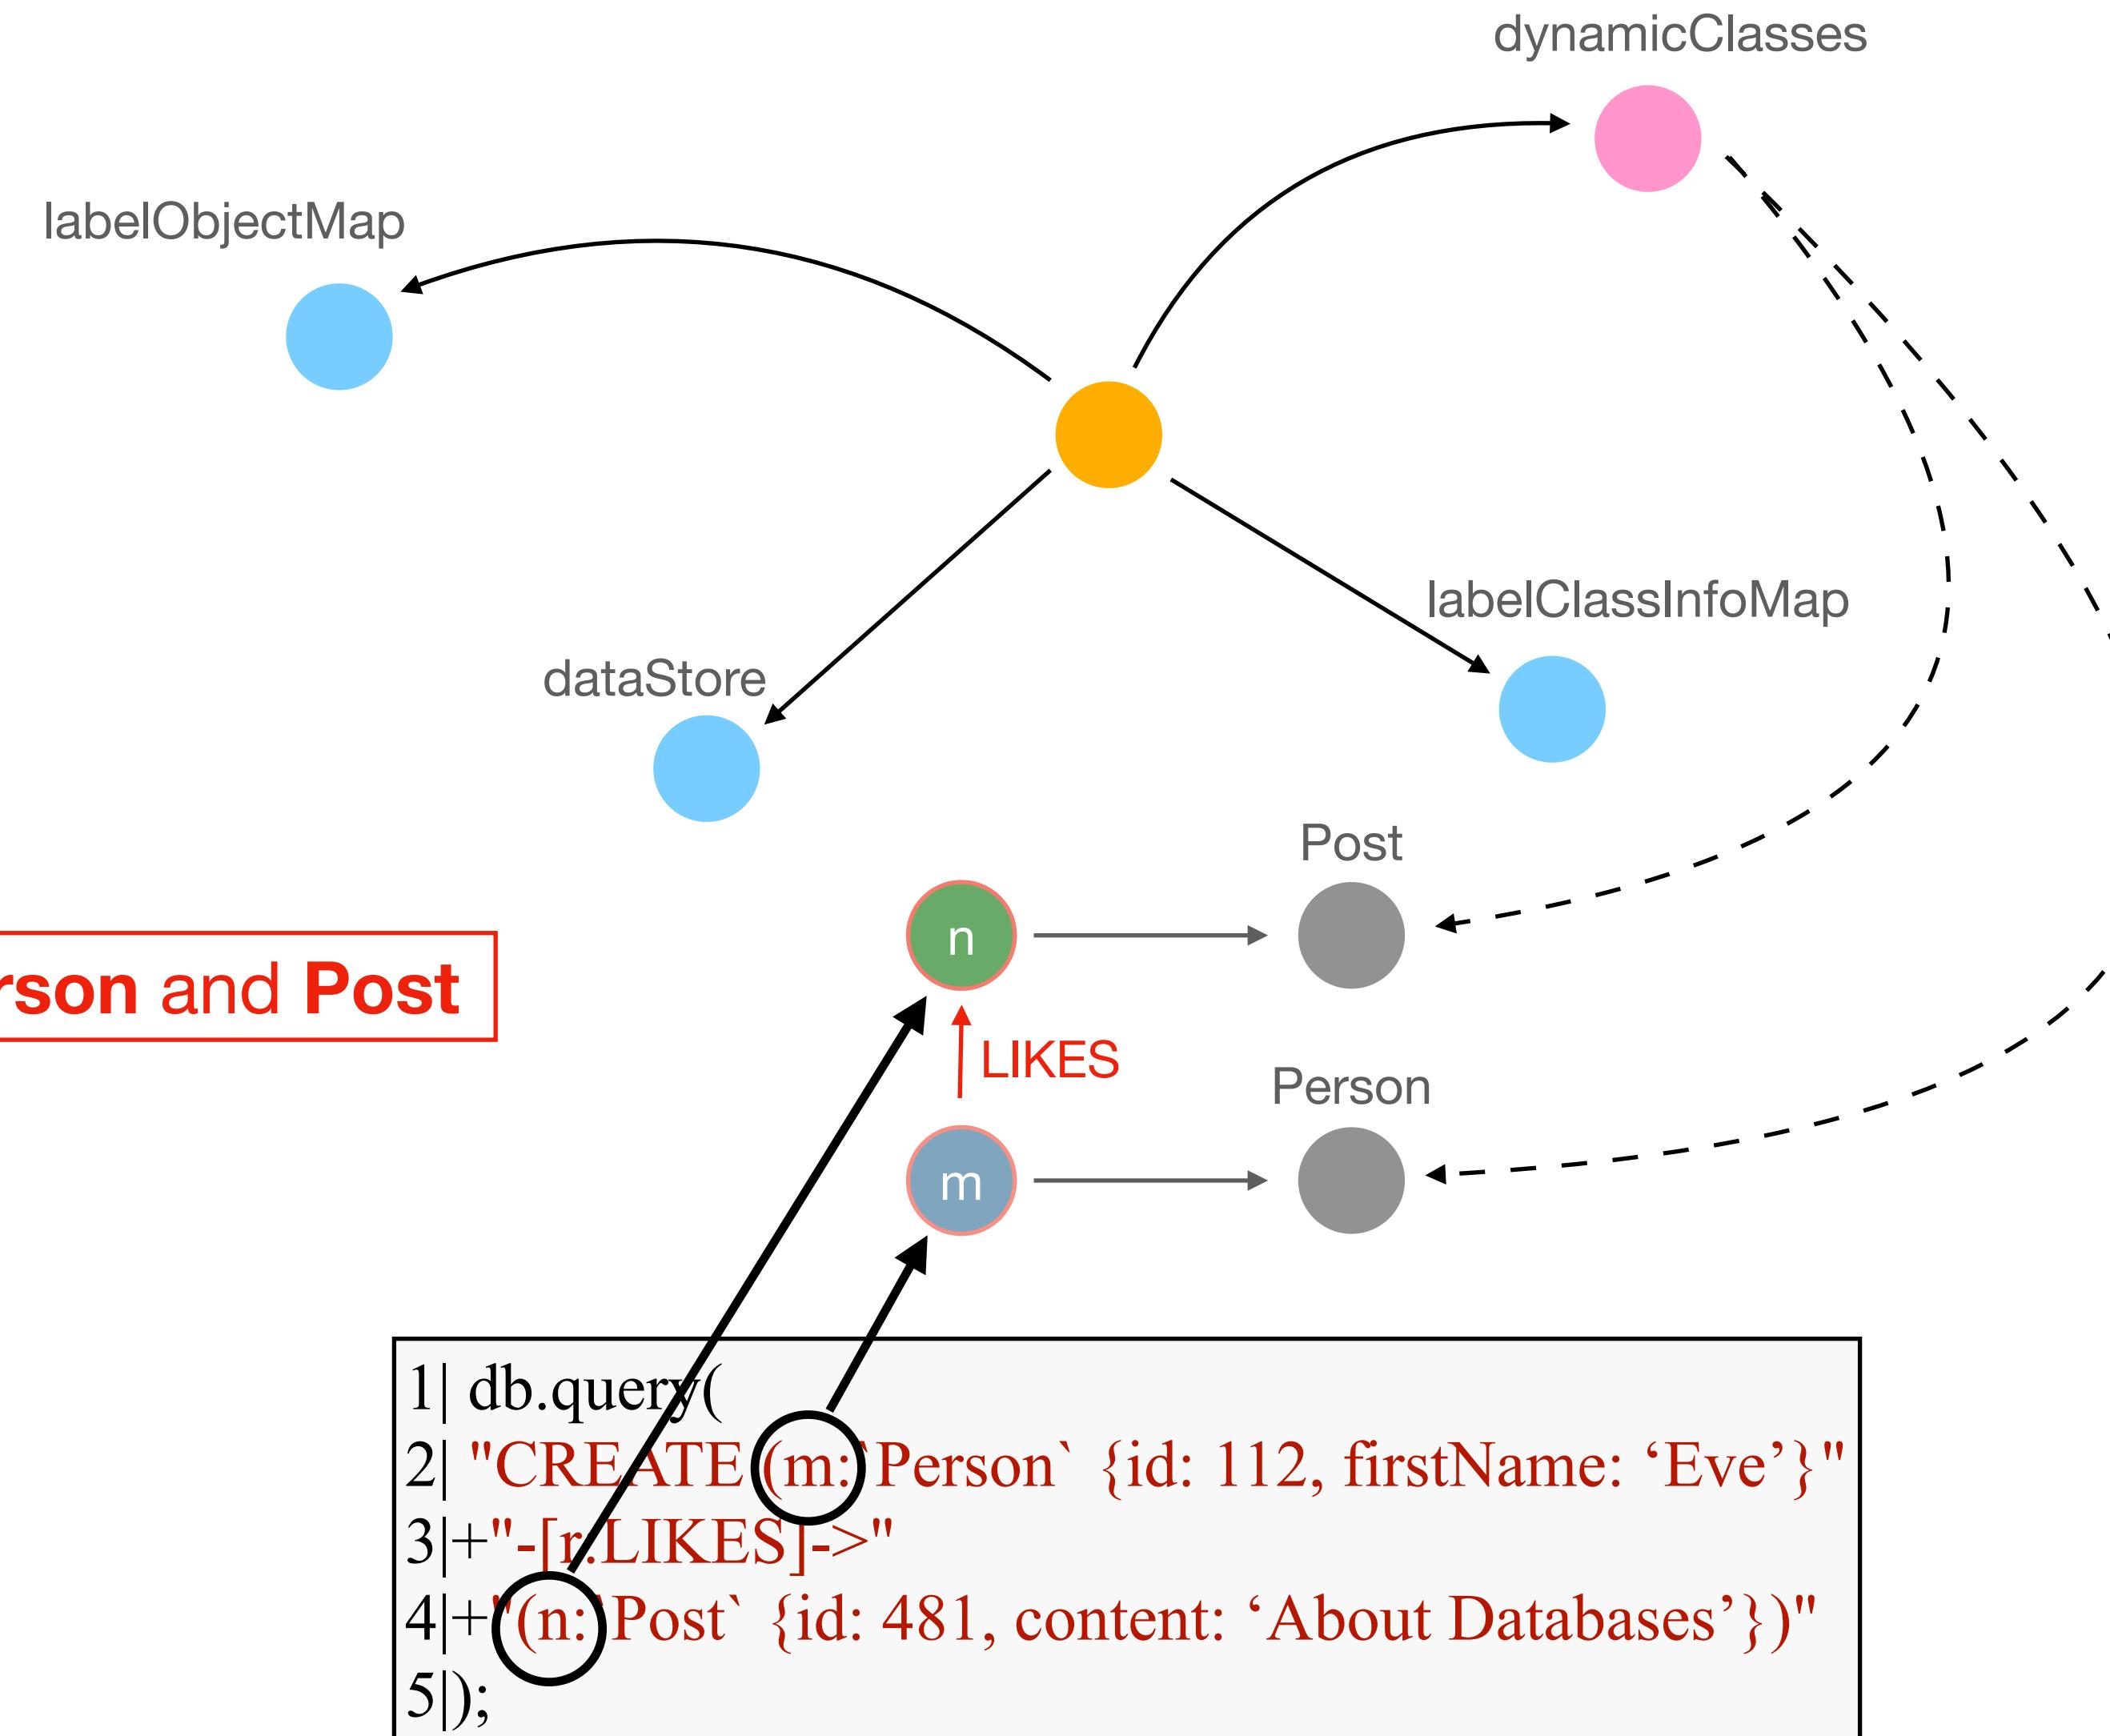
1. Runtime class creation for **Person** and **Post** using **ASM**
2. Store created classes in **dynamicClasses**
3. Cache class field names and types in **labelClassInfoMap** (not shown)



```
1| db.query(  
2| "CREATE (m:`Person` {id: 112, firstName: 'Eve'})"  
3| +"-[r:LIKES]->"  
4| +"(n:`Post` {id: 481, content: 'About Databases'})"  
5|);
```

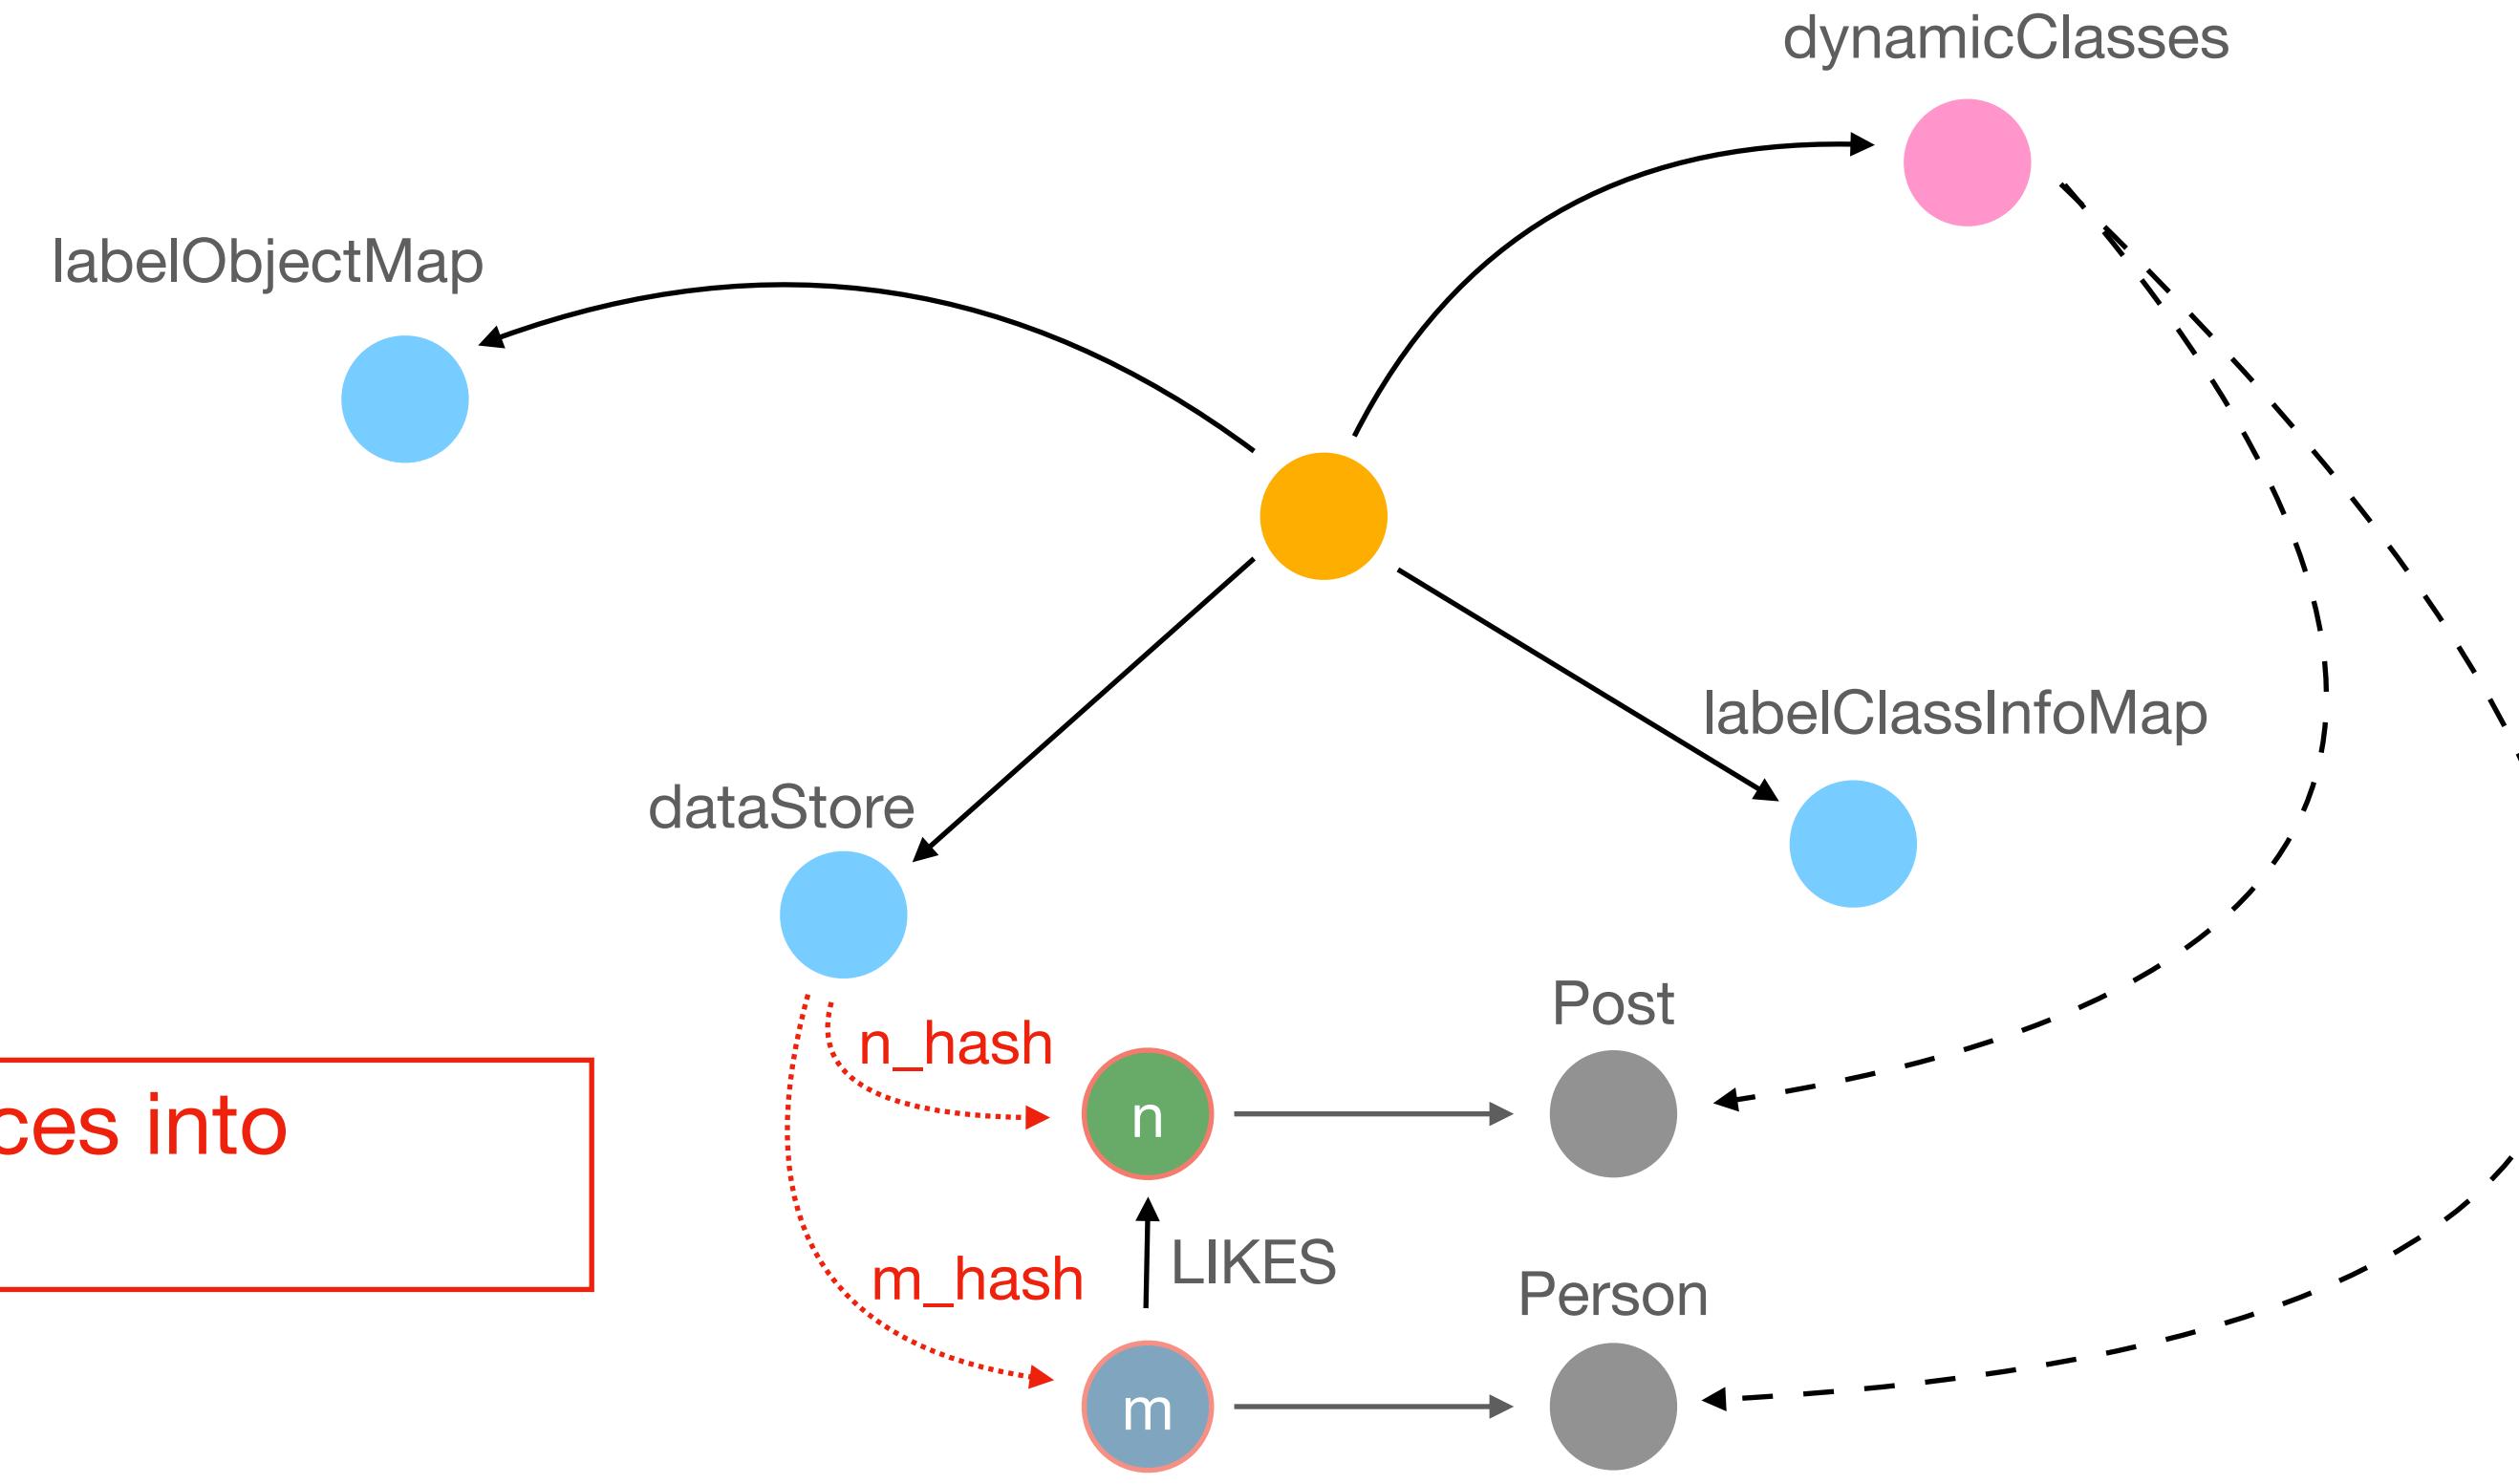


Example | Adding data without pre-defined schema



Example | Adding data without pre-defined schema

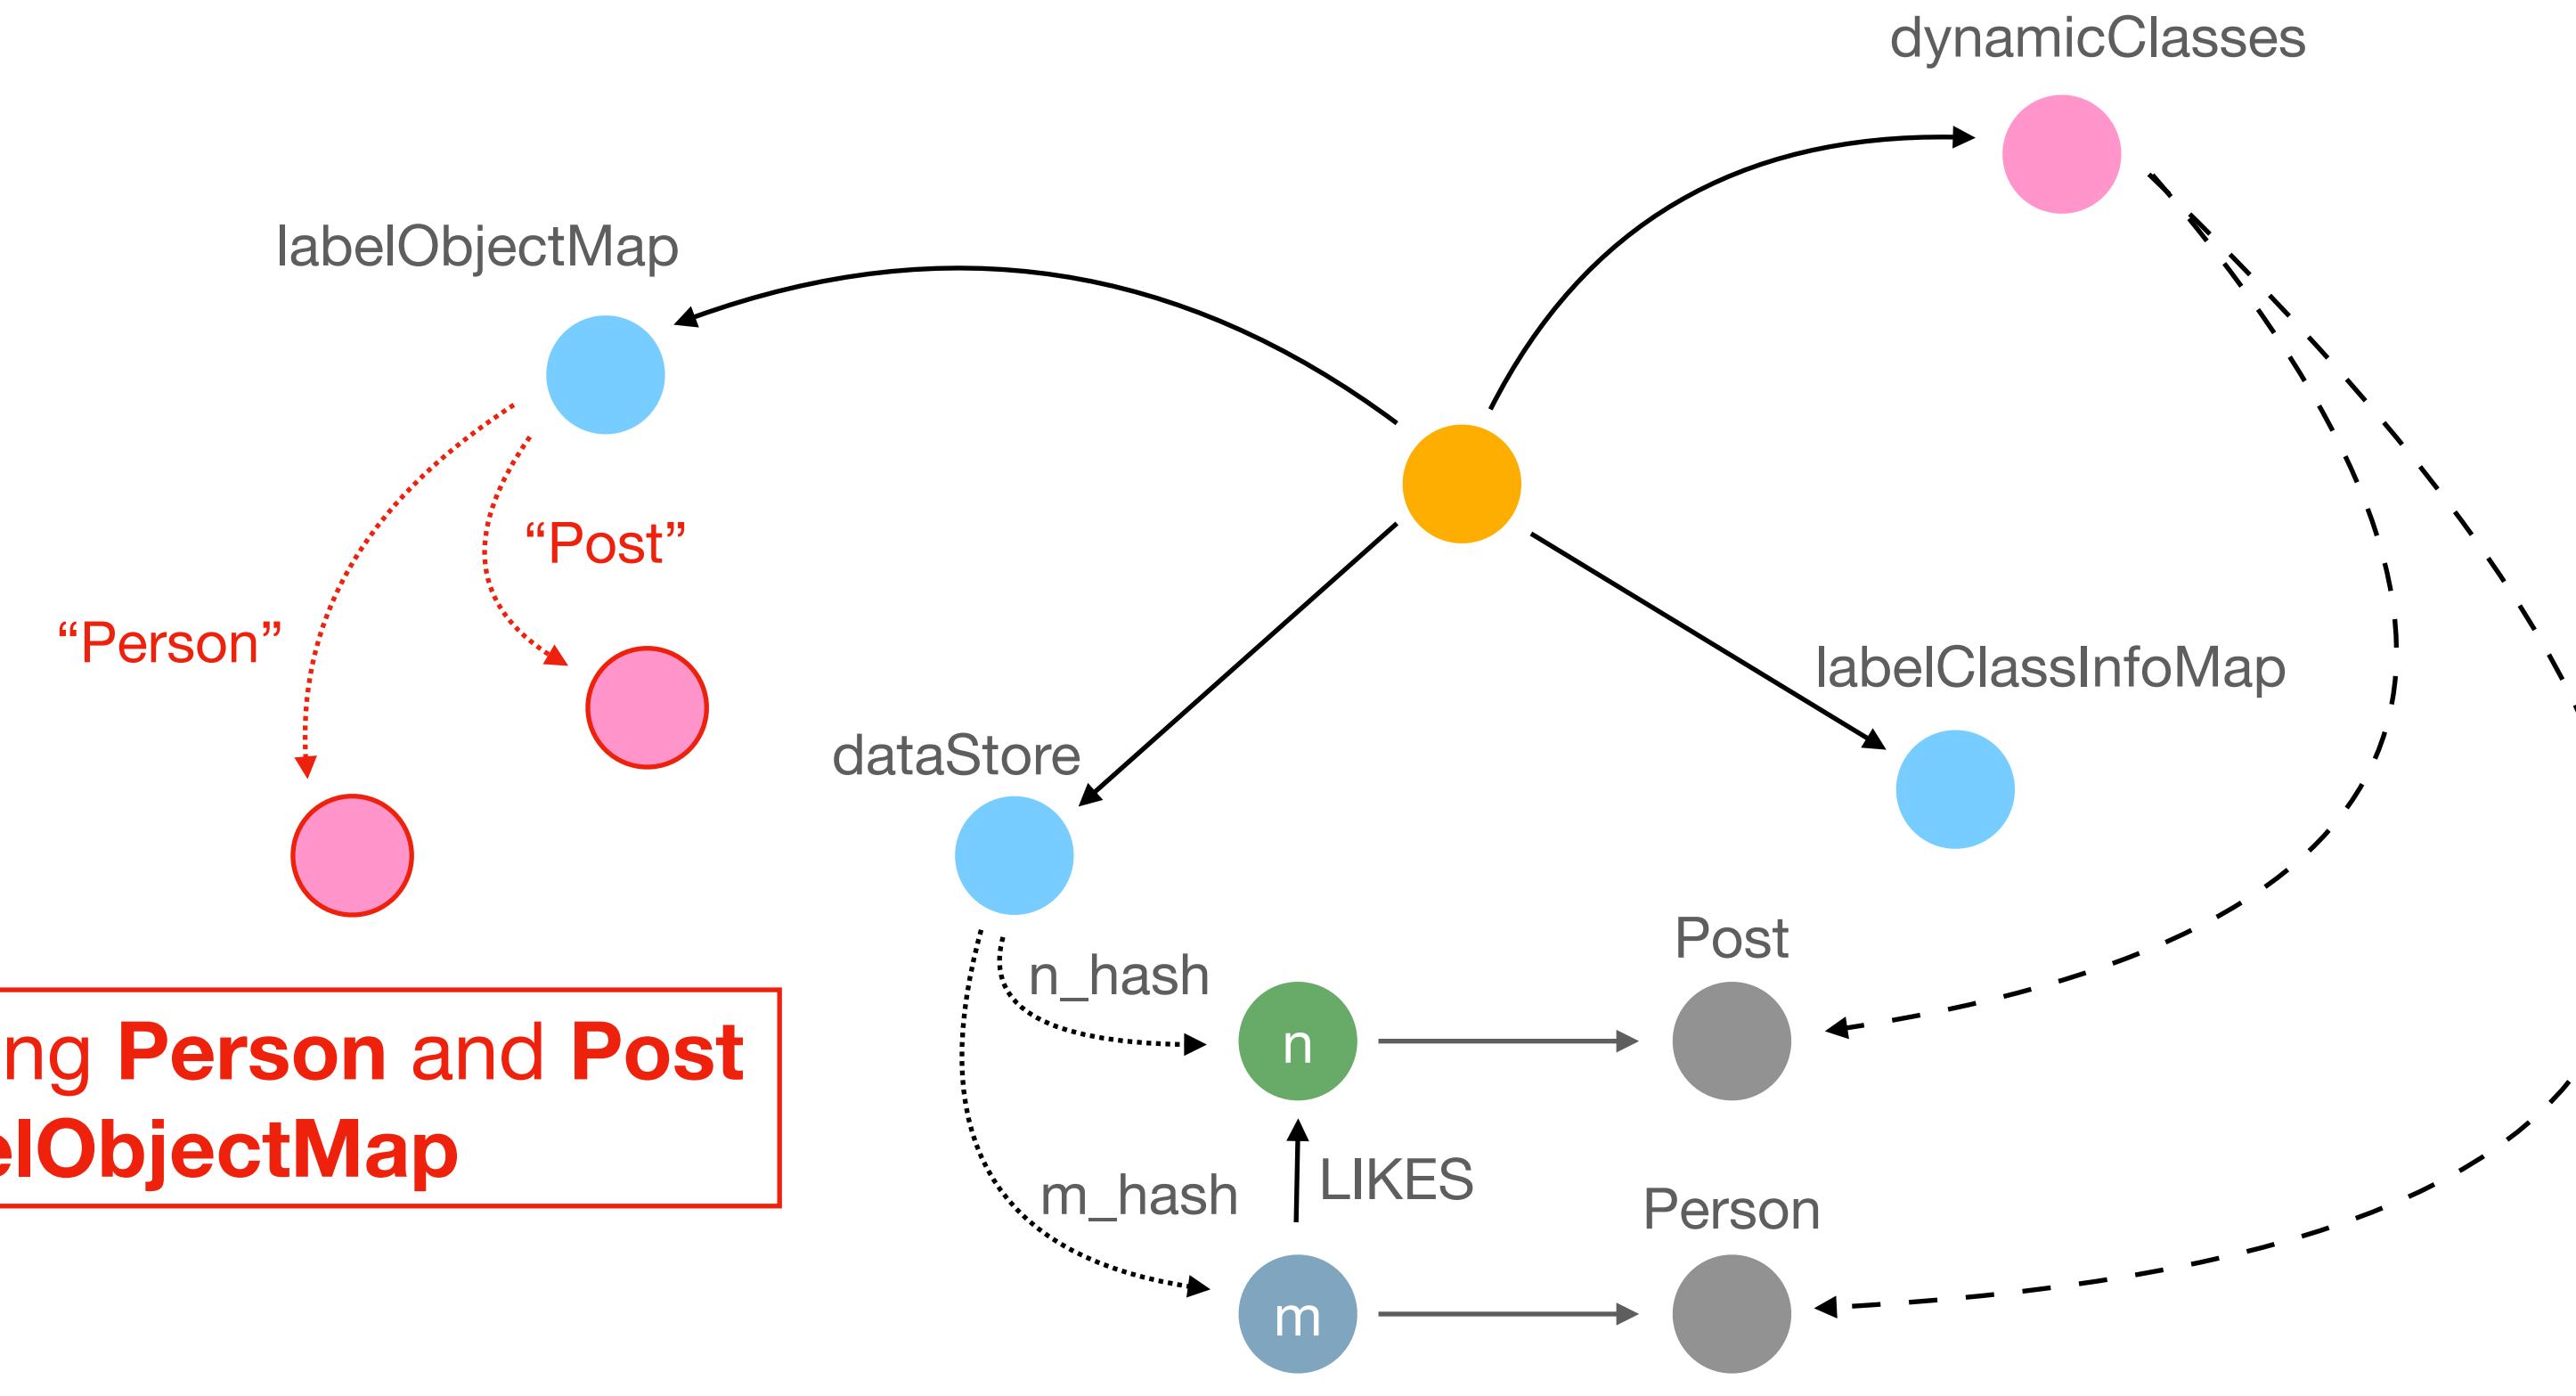
1. Capture created instances into datastore



```
1| db.query(  
2| "CREATE (m:'Person' {id: 112, firstName: 'Eve'})"  
3| +"-[r:LIKES]->"  
4| +"(n:'Post' {id: 481, content: 'About Databases'})"  
5|);
```



Example | Adding data without pre-defined schema

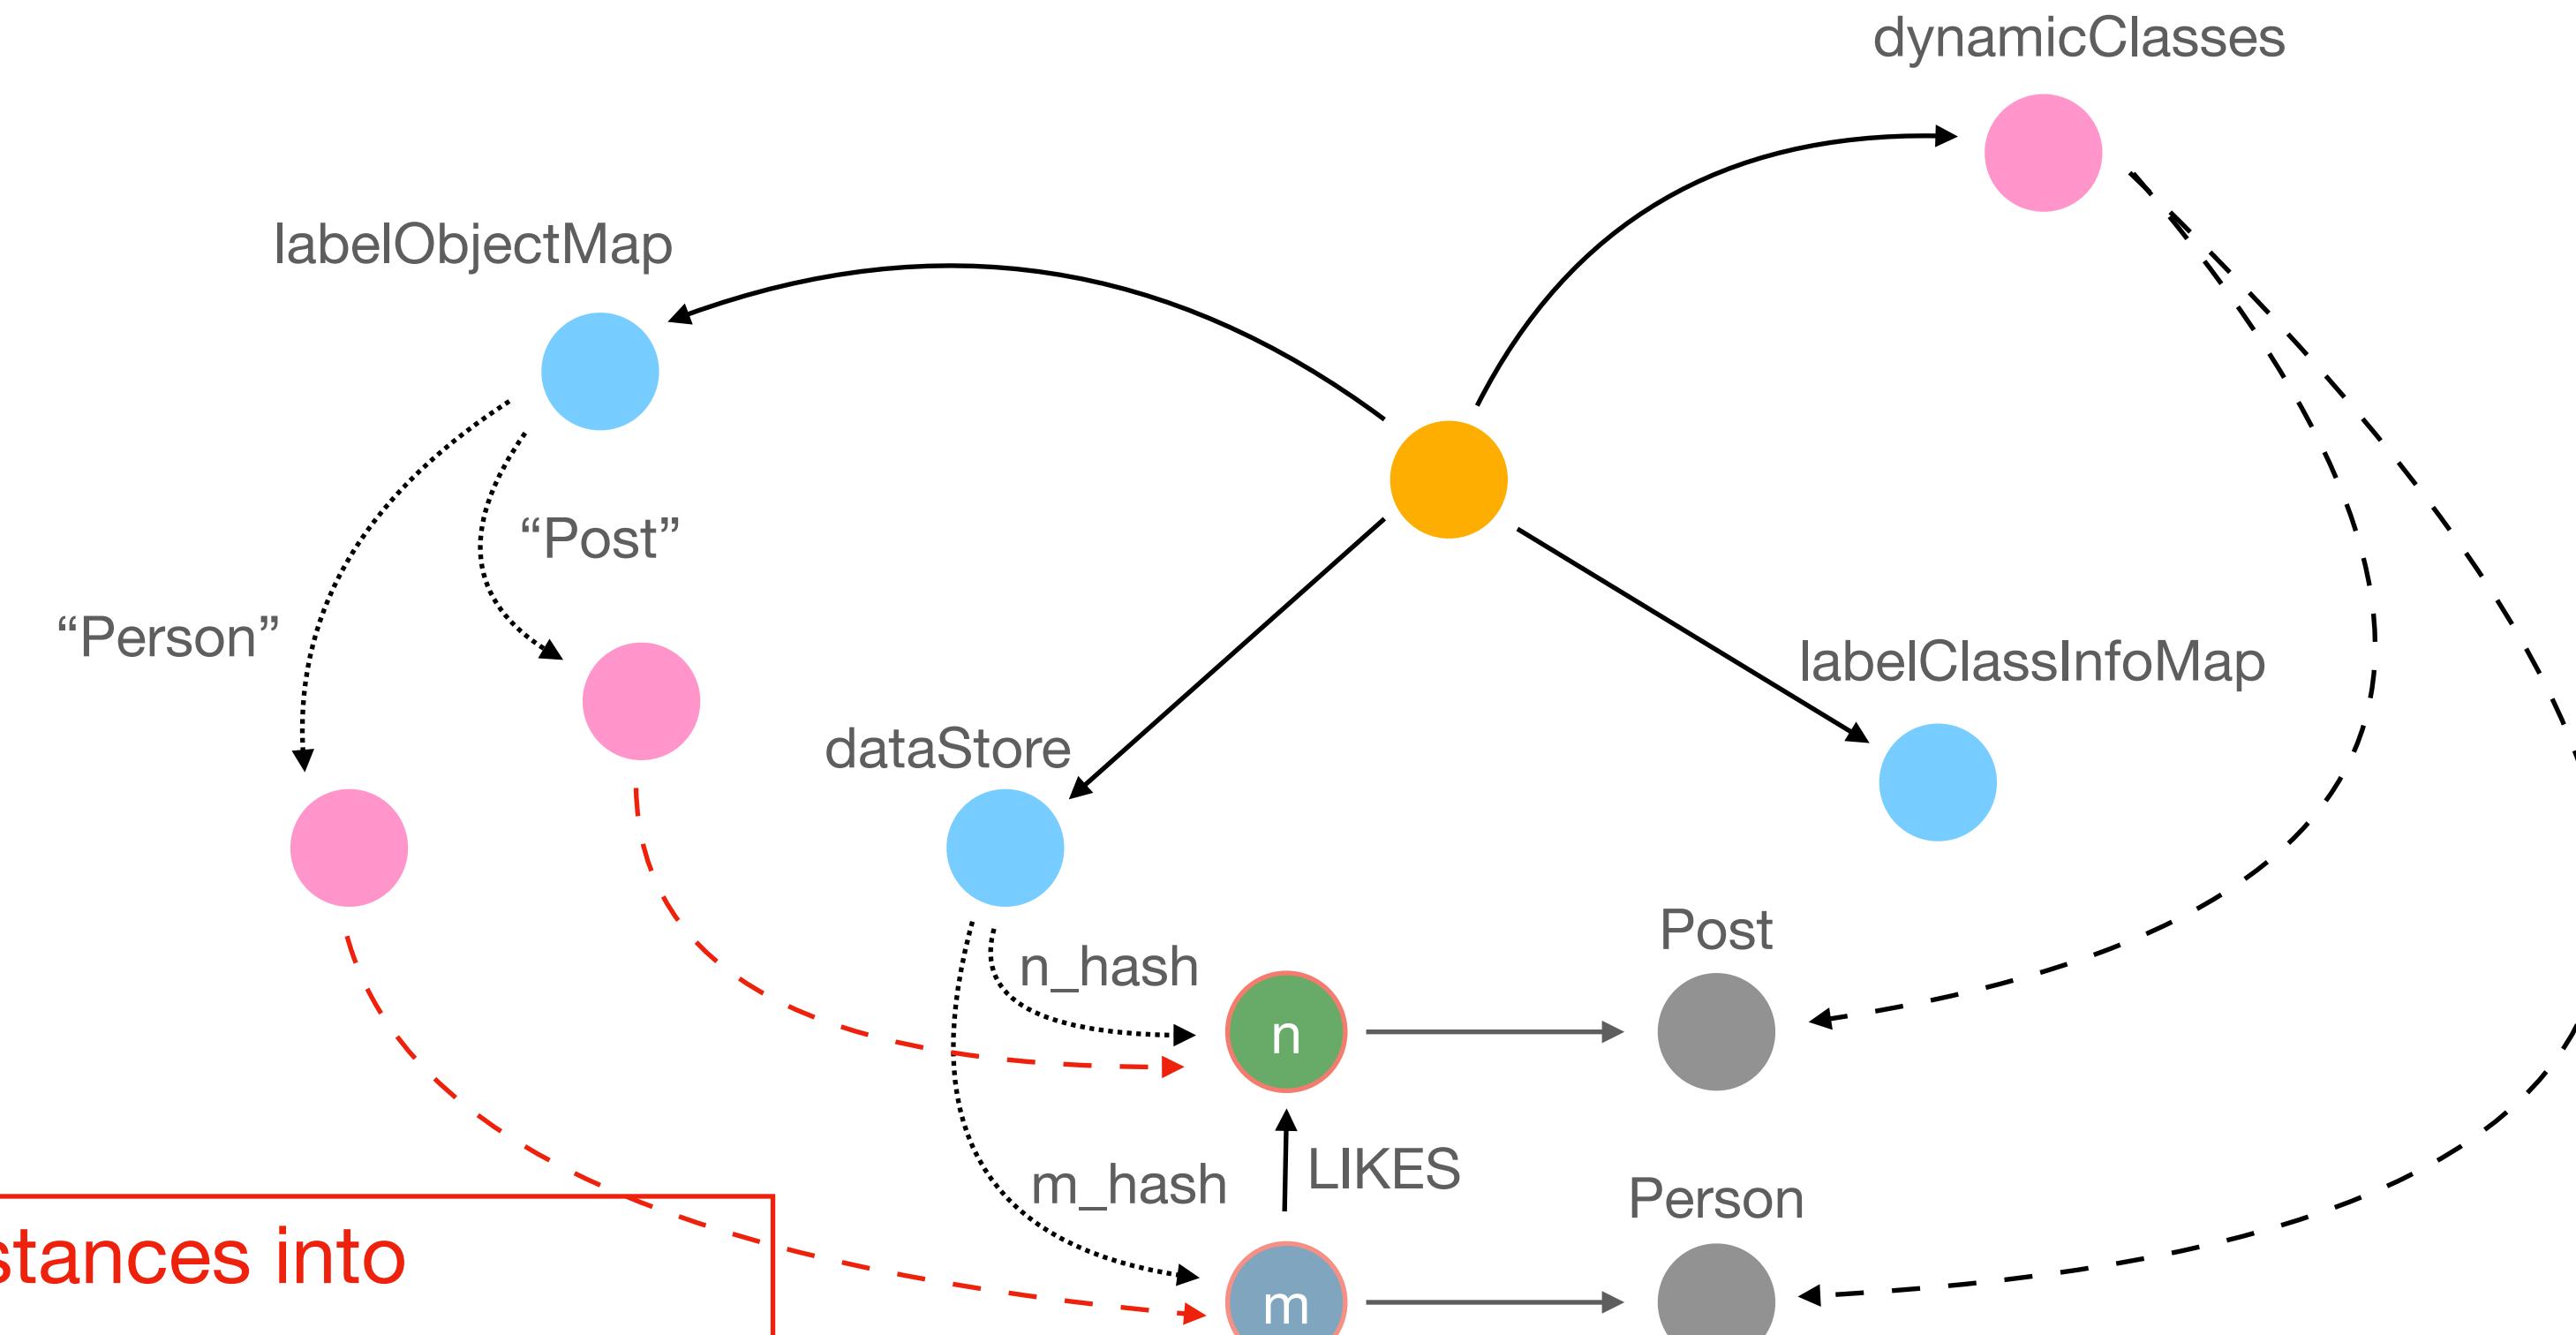


1. Add lists for storing **Person** and **Post** instances to **labelObjectMap**

```
1| db.query(  
2| "CREATE (m:'Person' {id: 112, firstName: 'Eve'})"  
3| +"-[r:LIKES]->"  
4| +(n:'Post' {id: 481, content: 'About Databases'})"  
5|);
```



Example | Adding data without pre-defined schema

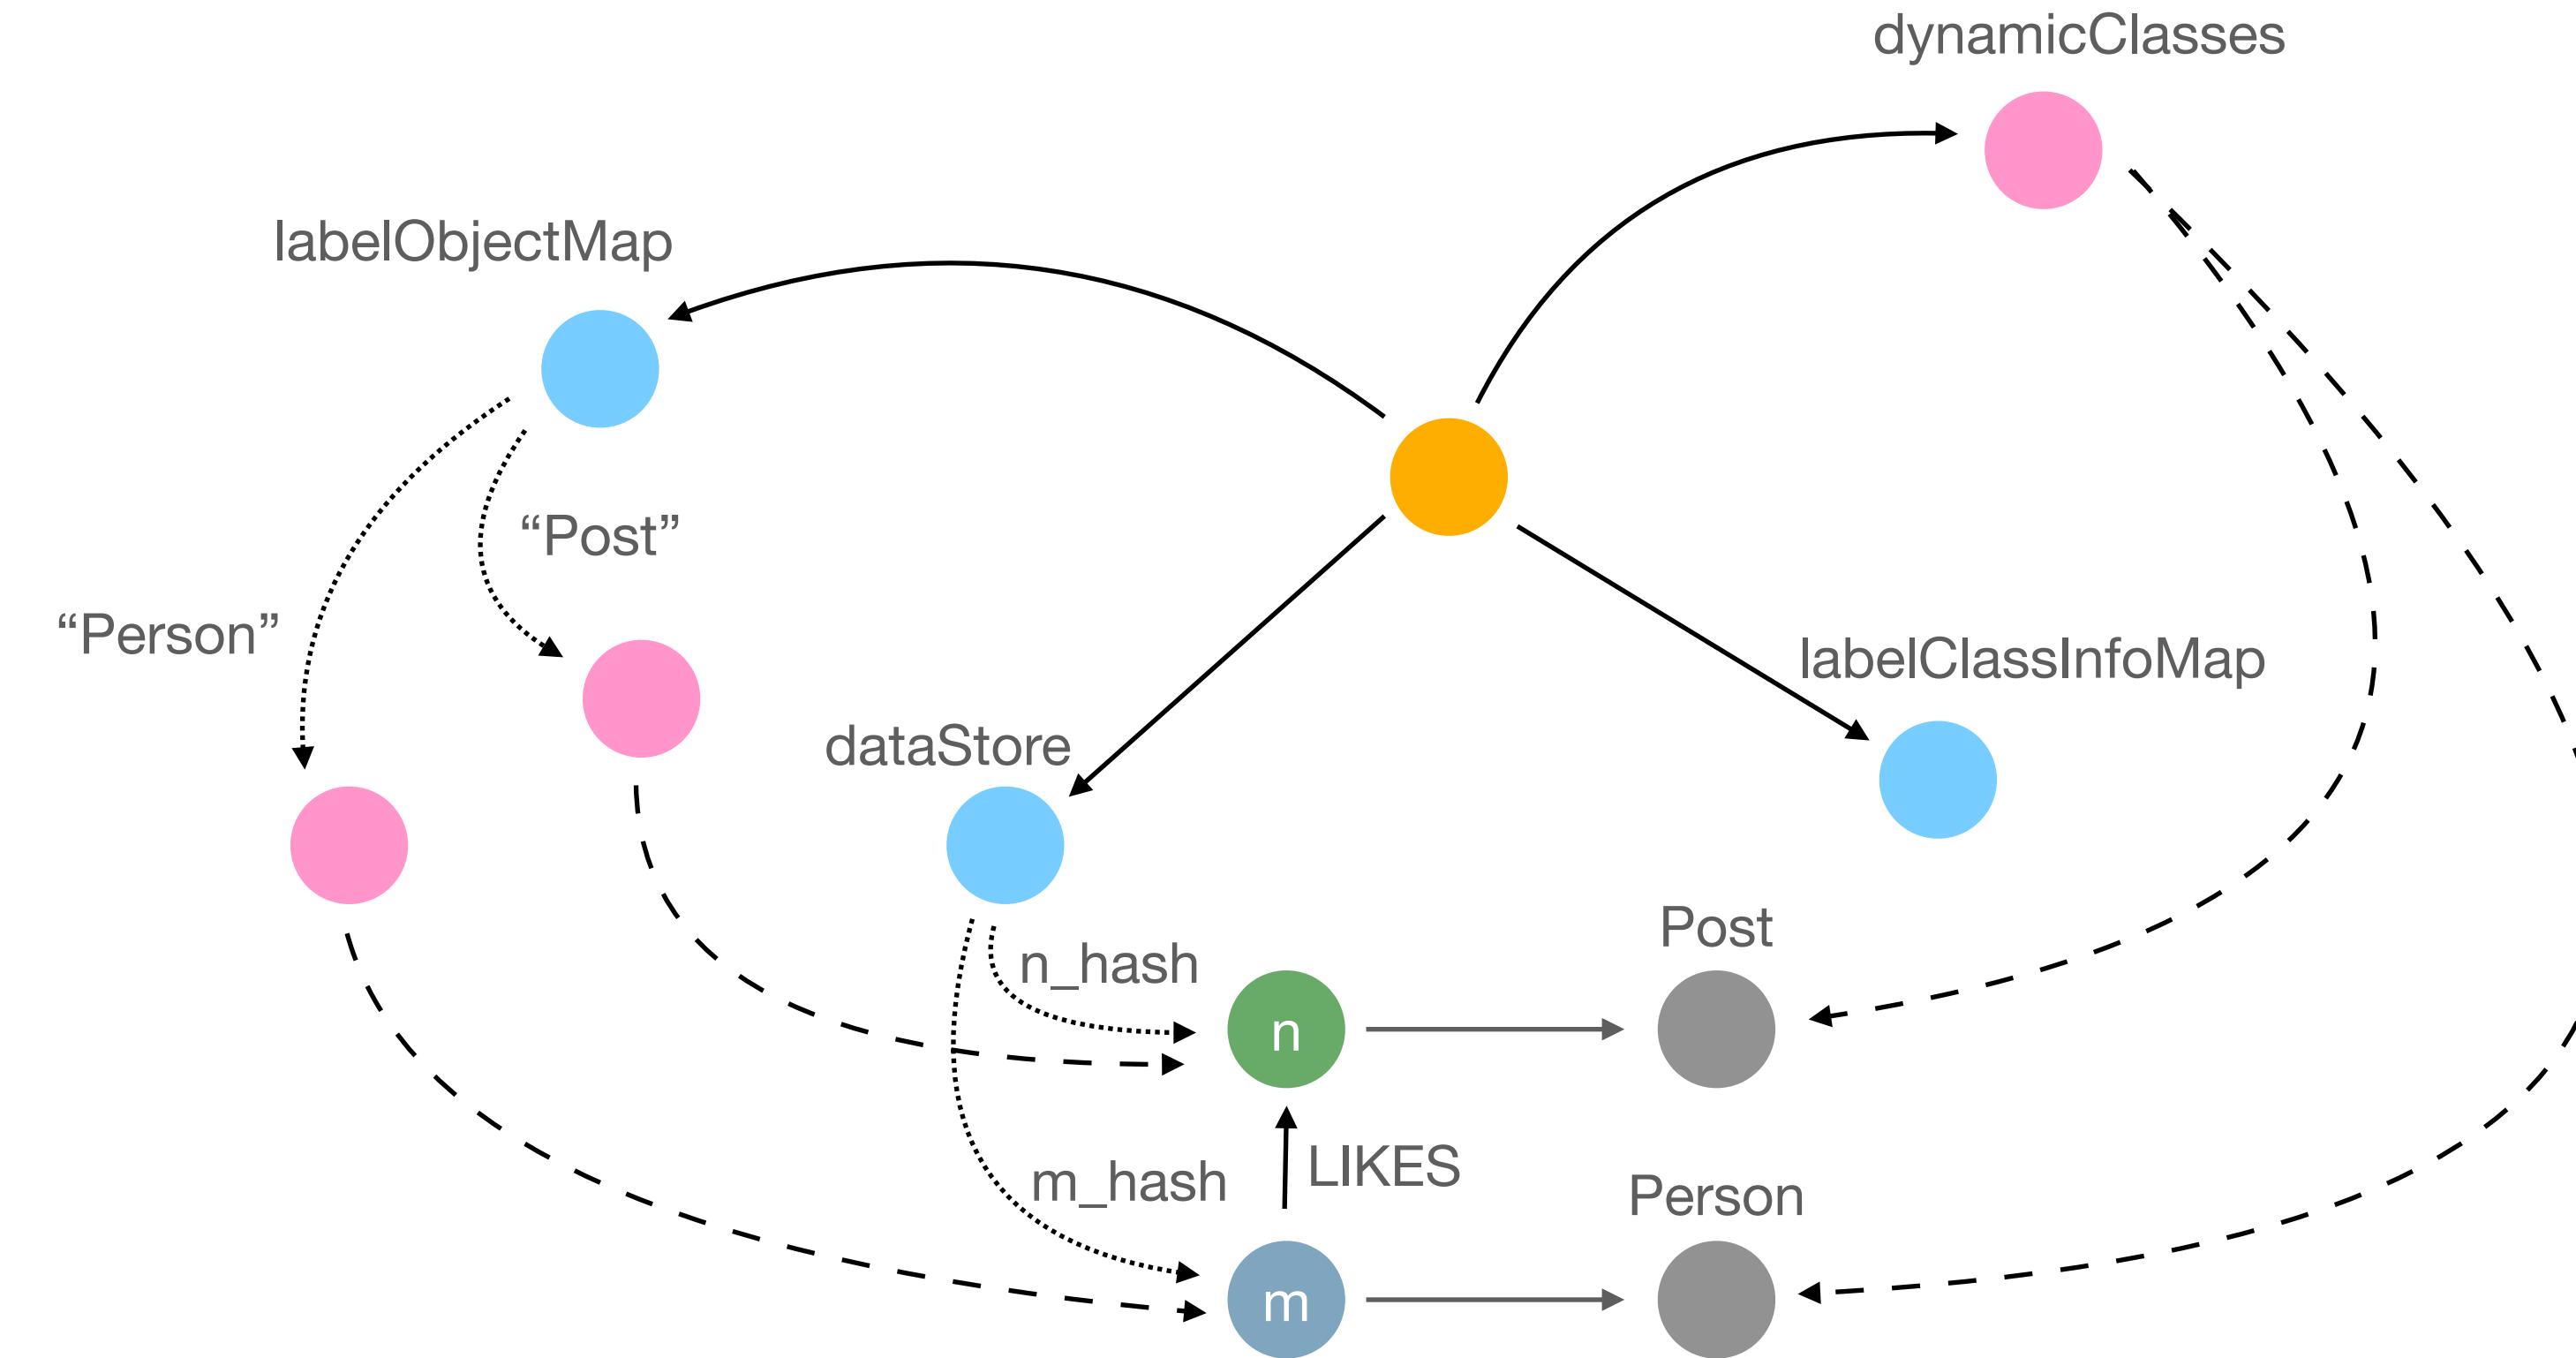


```
1| db.query(  
2| "CREATE (m:'Person' {id: 112, firstName: 'Eve'})"  
3| +"-[r:LIKES]->"  
4| +"(n:'Post' {id: 481, content: 'About Databases'})"  
5|);
```

L
E
G
E
N
D

	Class instance
	ArrayList instance
	HashMap instance
	eStore instance
	Post instance
	Person instance
	instanceof relation
	Key
	List element relation

Example | Deleting edges

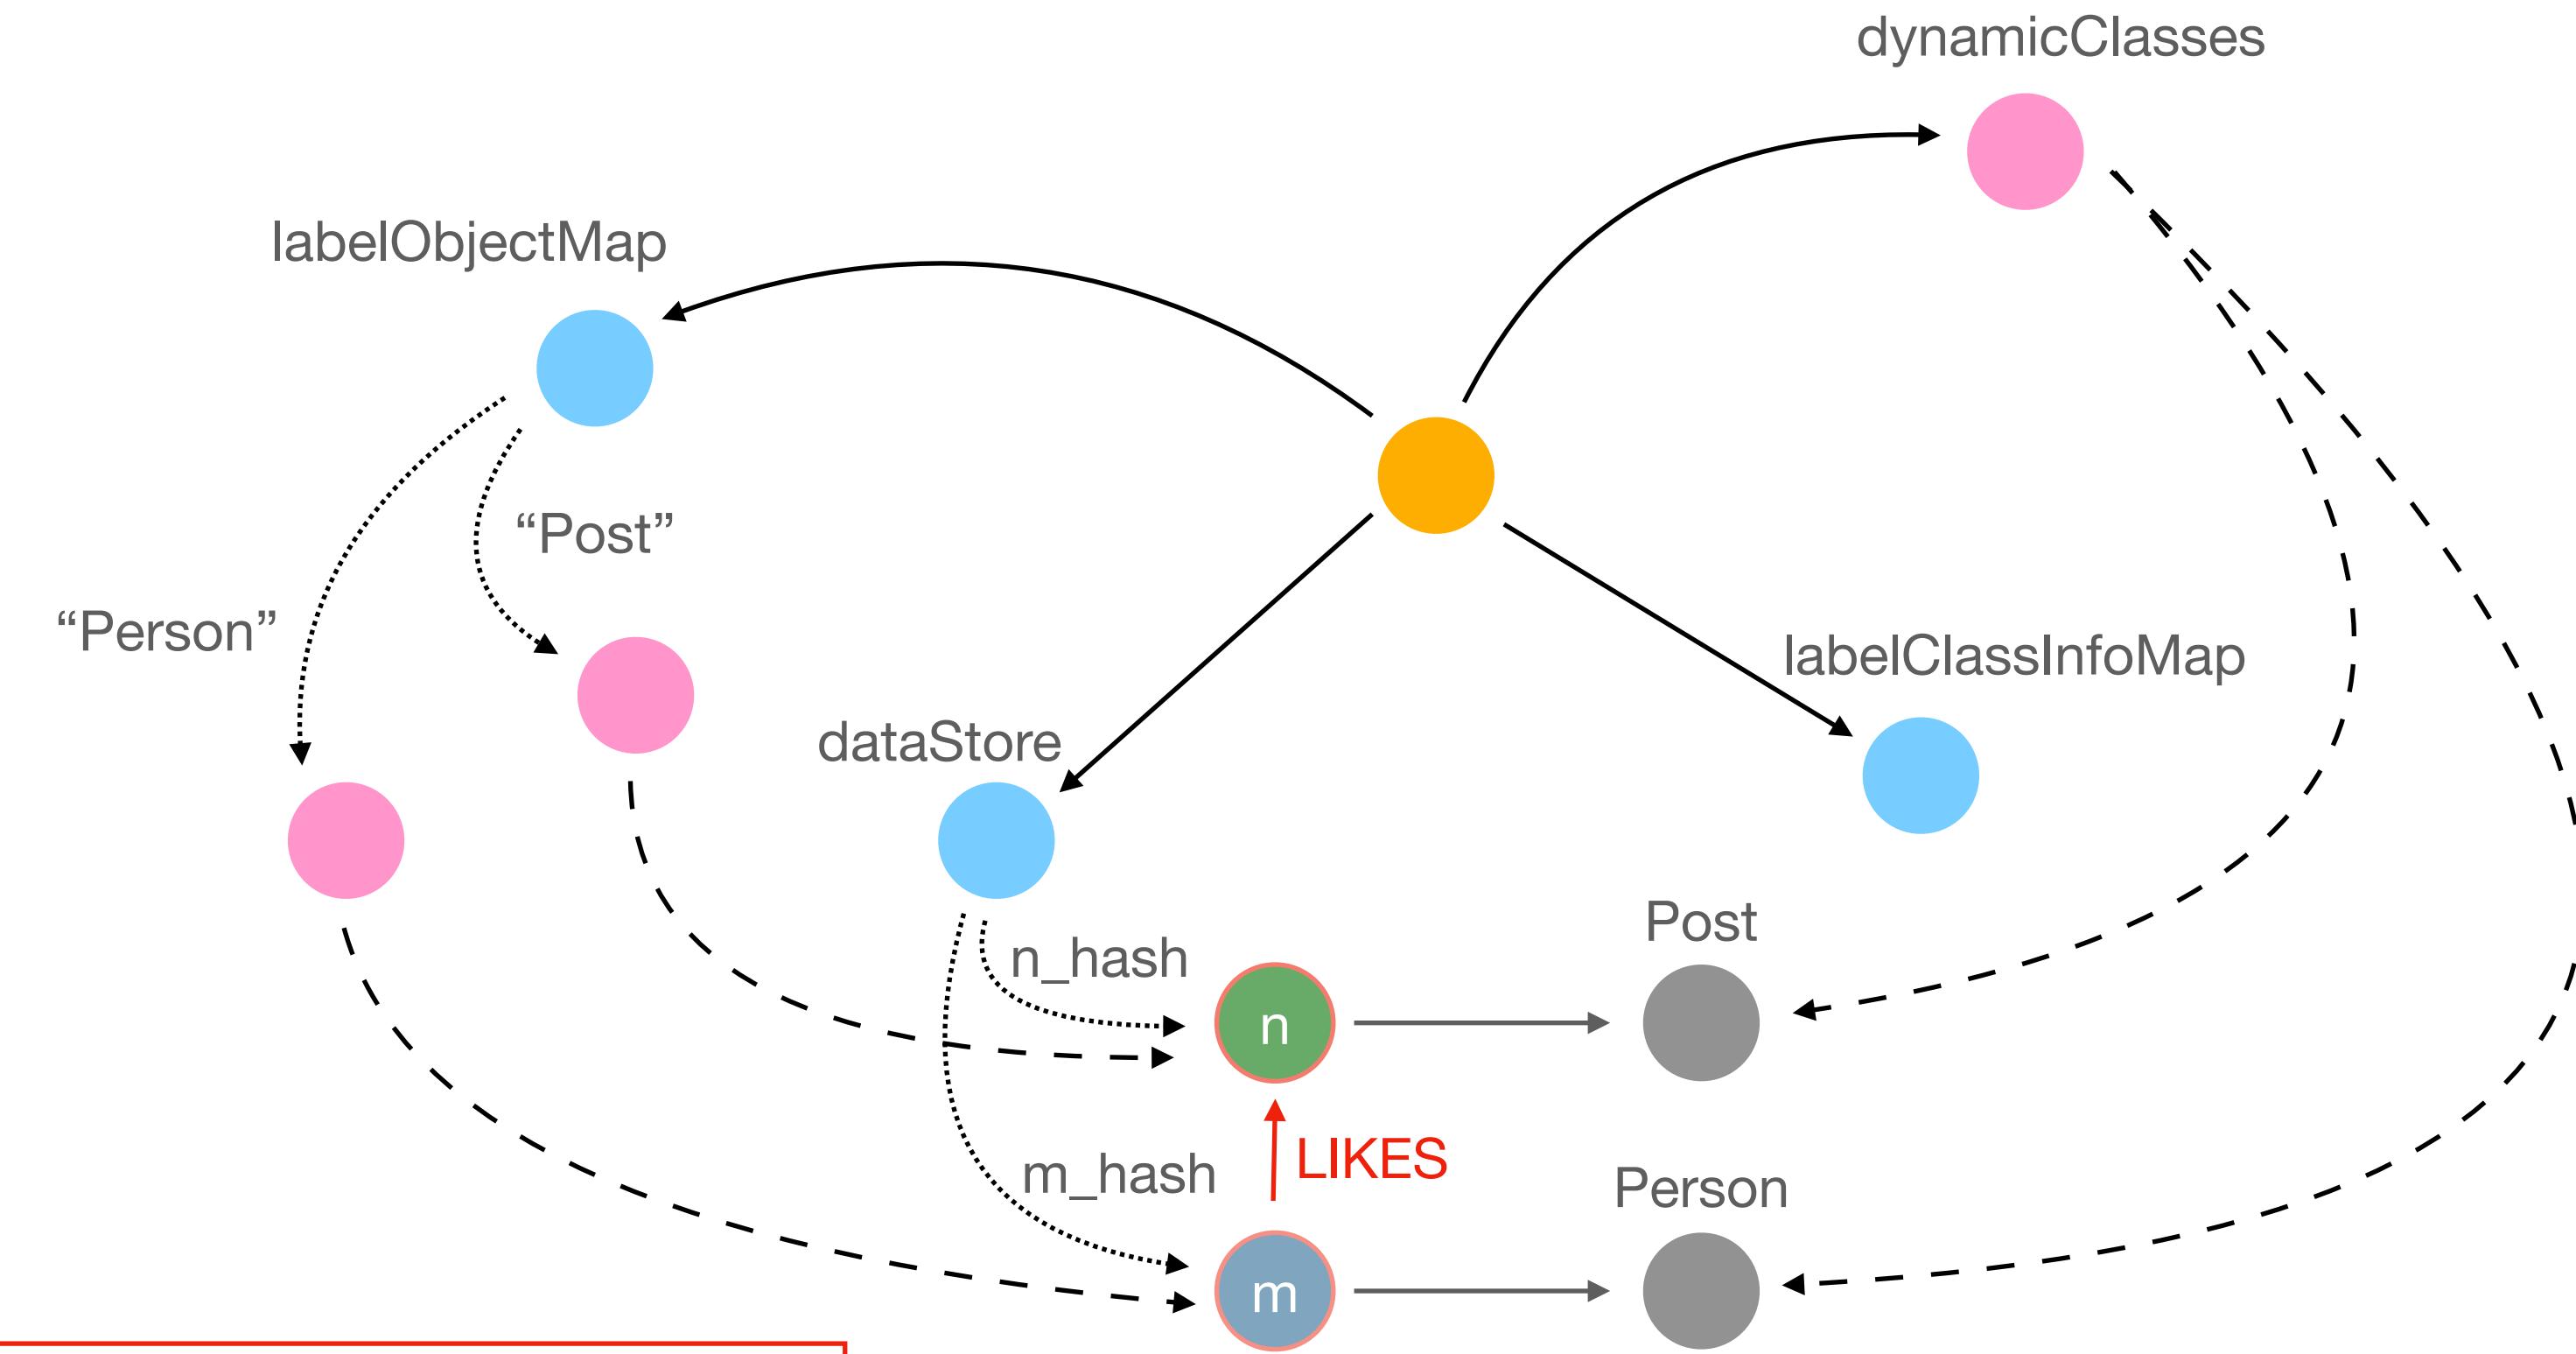


```
1| db.query(  
2| "MATCH (m:`Person` {id: 112})"  
3| +"-[r:LIKES]->"  
4| +"(n:`Post` {id: 481}) "  
5| +"DELETE r"  
6| );
```

L
E
G
E
N
D

	Class instance
	ArrayList instance
	HashMap instance
	eStore instance
	Post instance
	Person instance
	instanceof relation
	Key
	List element relation

Example | Deleting edges



1. Match the pattern

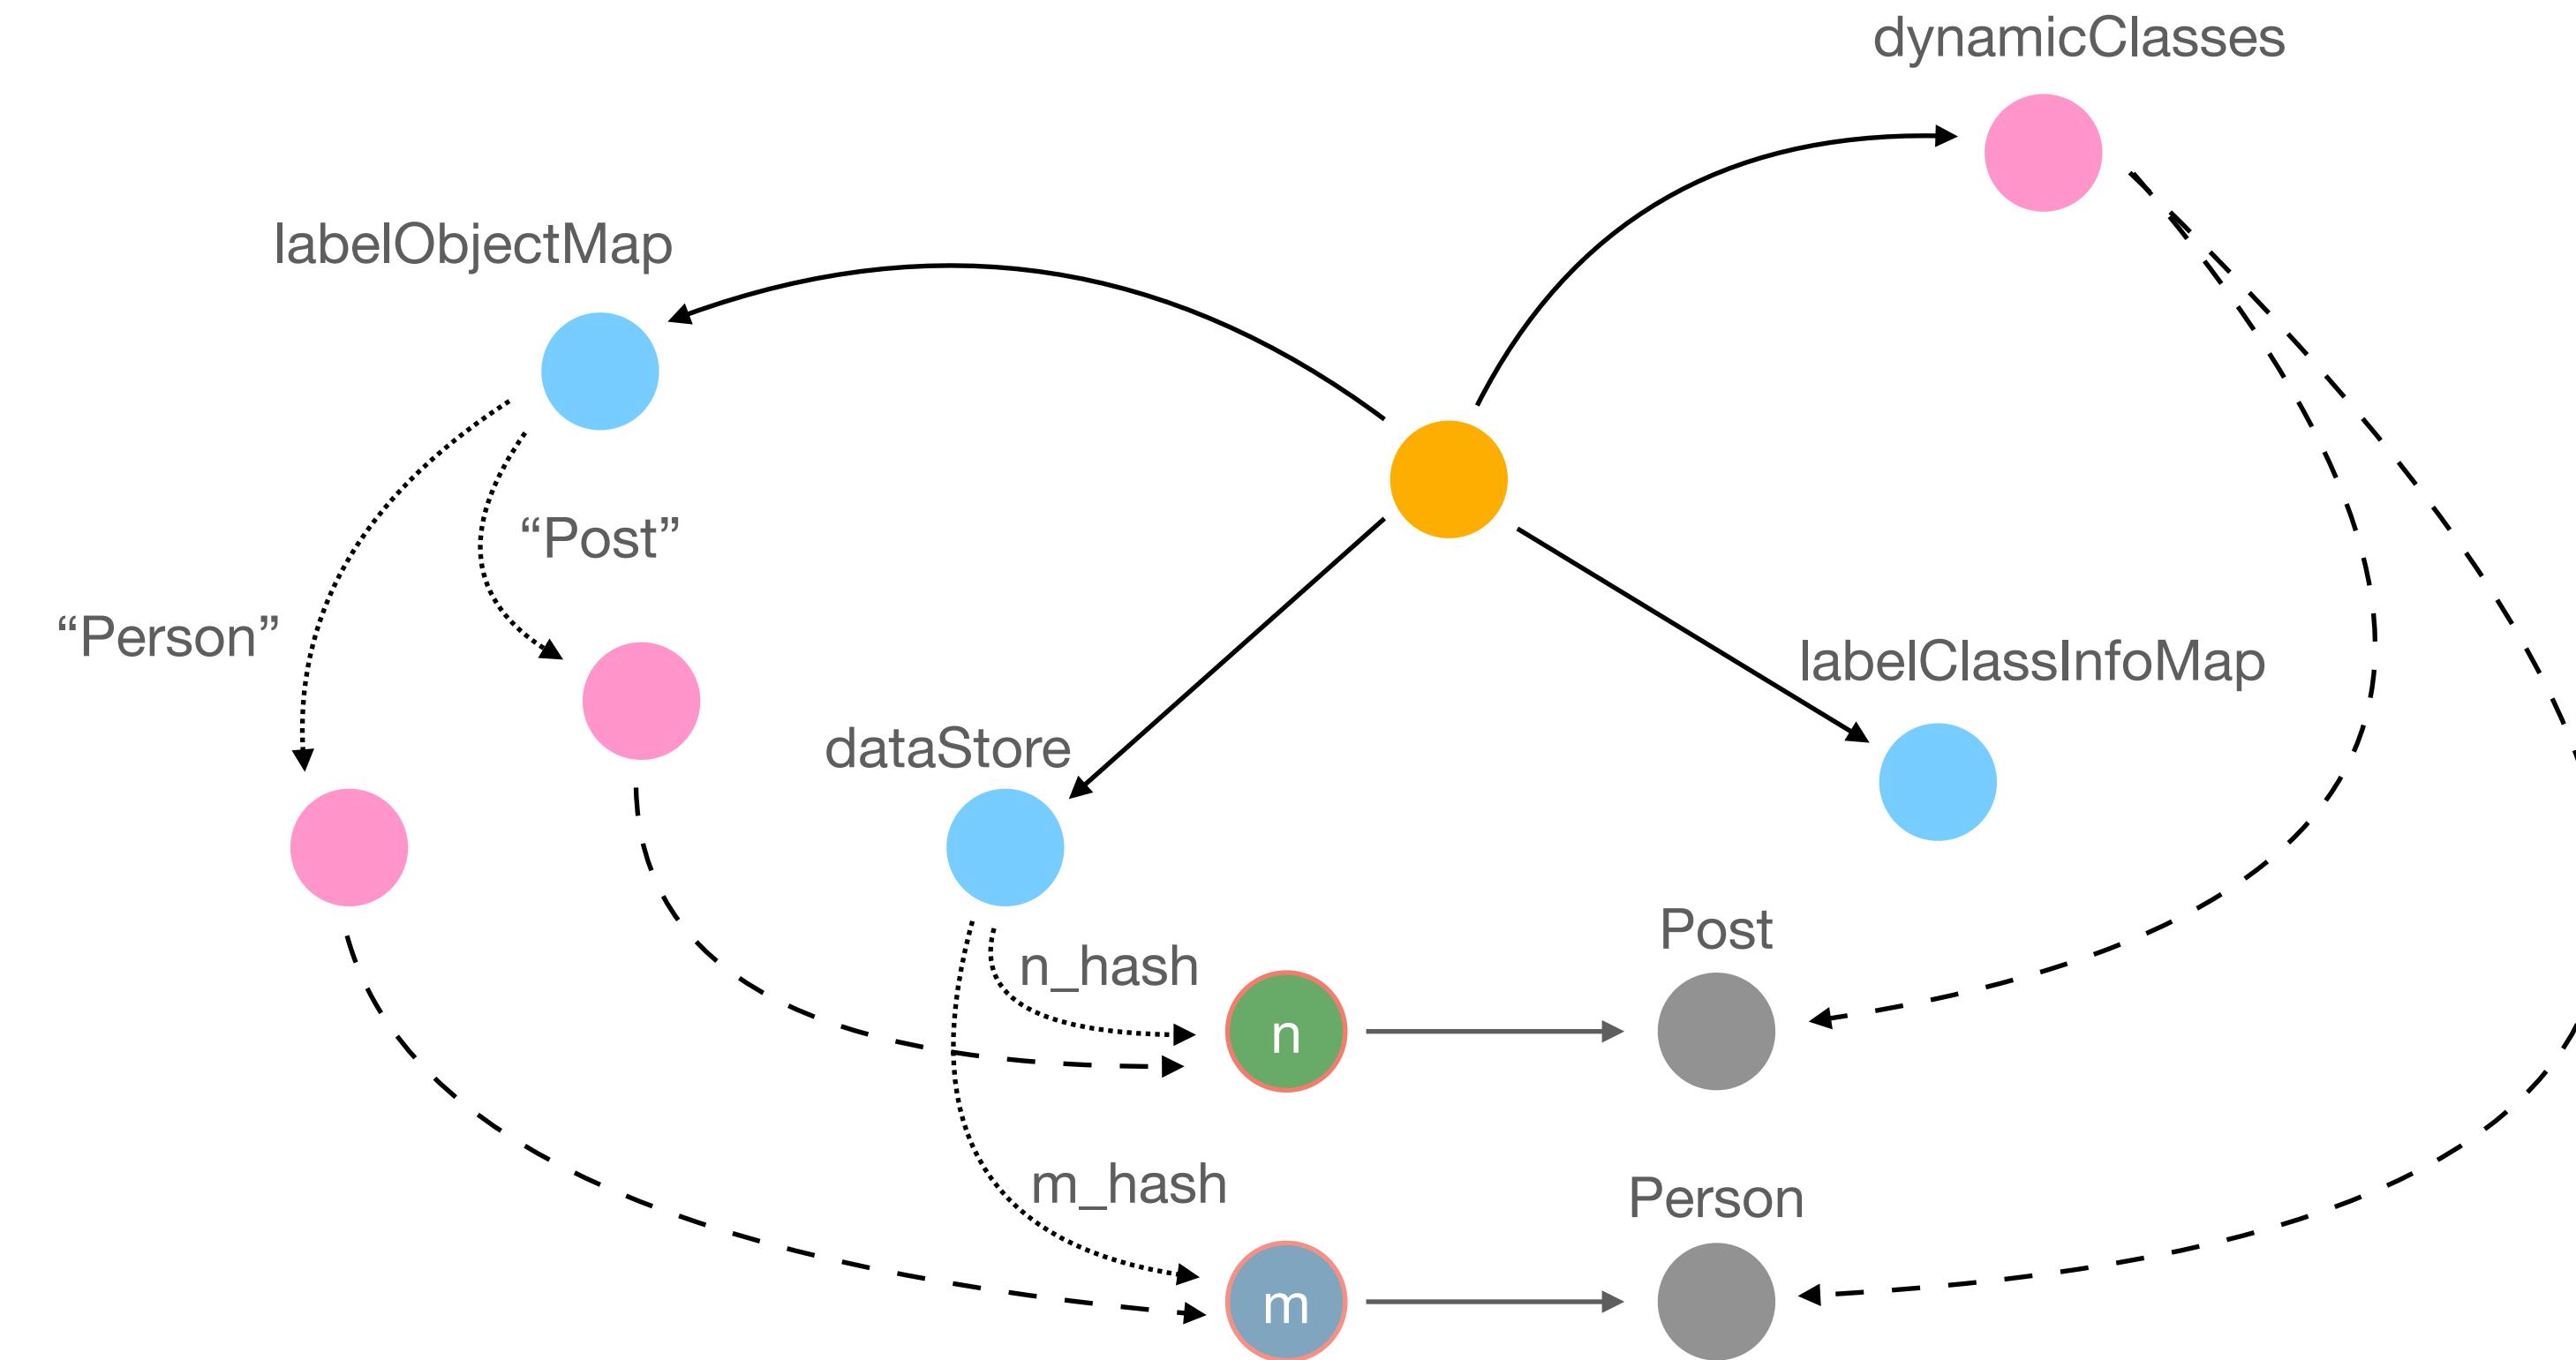
`(m:`Person` ...)-[r:LIKES]->(n:`Post` ...)`

```
1| db.query(  
2| "MATCH (m:`Person` {id: 112})"  
3| +"-[r:LIKES]->"  
4| +(n:`Post` {id: 481}) "  
5| +"DELETE r"  
6| );
```

L
E
G
E
N
D

Class instance	
ArrayList instance	
HashMap instance	
eStore instance	
Post instance	
Person instance	
instanceof relation	
Key	
HashMap instance	
List element relation	

Example | Deleting edges



1. Delete edge LIKES (r)

```
1| db.query(  
2| "MATCH (m:`Person` {id: 112})"  
3| +"-[r:LIKES]->"  
4| +"(n:`Post` {id: 481}) "  
5| +"DELETE r"  
6| );
```



Results | LDBC SNB Benchmark

- Benchmark setup:** We compare end-to-end and component latencies for Neo4j (server & in-memory) versus eStore on 9 LDBC SNB queries across scale factors 0.1–10.

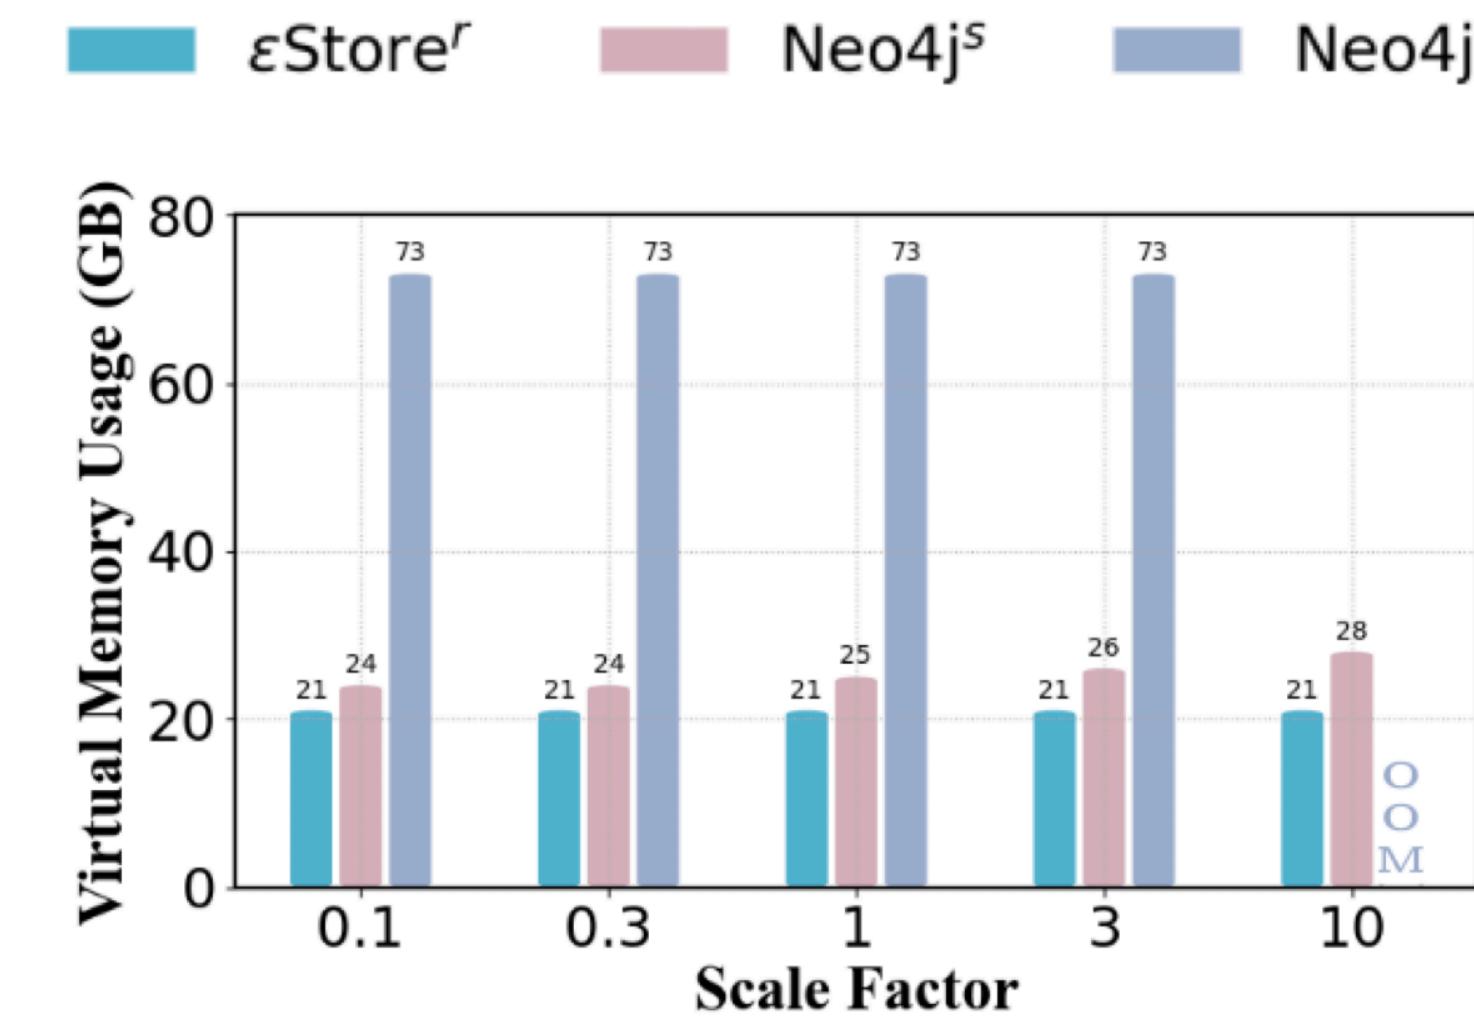
- eStore completes every run within the 63 GB cap and delivers 5–13 × lower latency than Neo4j, while Neo4jⁱ hits OOM on all scale factor 10 queries.

Query	SF	Neo4j ^s					Neo4j ⁱ					ϵ Store ^r				
		t_{Pa}	t_{Pl}	t_{Ex}	t_{Bd}	T_{tot}	t_{Pa}	t_{Pl}	t_{Ex}	t_{Bd}	T_{tot}	t_{Pa}	t_{Pl}	t_{Ex}	t_{Bd}	T_{tot}
Q_{SNB}^A	0.1	11	29	0	40	46	1	33	3	37	38	0	0	0	0	0
	0.3	11	28	0	39	44	1	33	5	39	40	0	0	1	1	3
	1	10	24	0	34	38	0	31	24	55	55	0	0	0	1	1
	3	12	24	0	36	40	2	32	40	75	76	0	0	1	4	5
	10	11	23	0	35	39			OOM			0	0	0	4	5
Q_{SNB}^B	0.1	10	24	0	34	39	0	32	2	35	36	0	0	0	0	0
	0.3	12	27	0	39	43	0	33	5	38	39	0	0	0	1	2
	1	11	26	0	37	41	1	34	9	44	45	0	0	0	0	1
	3	10	23	0	33	37	0	33	52	86	87	0	0	1	4	6
	10	11	24	0	35	39			OOM			0	0	0	4	4
Q_{SNB}^C	0.1	11	25	0	36	41	0	32	2	34	35	0	0	0	1	1
	0.3	10	27	0	37	41	0	32	4	37	37	0	0	0	3	4
	1	10	24	0	34	38	0	30	18	49	50	0	0	0	3	3
	3	10	24	0	34	38	0	31	21	52	53	0	0	0	8	9
	10	11	25	0	36	41			OOM			0	0	0	19	19
Q_{SNB}^D	0.1	16	30	0	46	51	1	34	0	36	36	0	0	0	0	1
	0.3	15	30	0	45	50	1	34	0	35	36	0	0	1	2	4
	1	13	25	0	38	43	0	34	0	34	35	0	0	0	1	1
	3	16	26	0	43	47	3	34	0	37	38	0	0	1	2	3
	10	15	26	0	42	46			OOM			0	0	0	4	5
Q_{SNB}^E	0.1	10	23	0	33	38	1	30	0	31	31	0	0	0	7	7
	0.3	11	21	0	32	36	2	34	0	36	37	0	0	0	29	29
	1	11	22	0	34	38	4	29	0	33	33	0	0	0	114	115
	3	11	20	0	31	36	6	29	0	36	37	0	0	1	319	320
	10	11	20	0	31	35			OOM			0	0	0	866	866
Q_{SNB}^F	0.1	10	15	29	54	58	0	24	44	67	68	0	0	0	5	6
	0.3	10	16	66	92	97	1	25	107	133	134	0	0	1	15	16
	1	9	16	212	237	241	1	24	286	311	312	0	0	0	35	35
	3	8	16	530	554	559	0	23	754	777	778	0	0	0	87	87
	10	11	18	1531	1559	1564			OOM			0	0	0	258	258
Q_{SNB}^G	0.1	9	16	31	57	61	2	26	59	88	88	0	0	0	6	6
	0.3	9	16	102	128	133	0	24	206	230	231	0	0	0	30	31
	1	10	15	423	448	453	1	24	728	753	754	0	0	0	74	75
	3	9	15	1298	1322	1326	3	25	2484	2511	2512	0	0	0	244	245
	10	10	15	4393	4418	4422			OOM			0	0	0	915	915
Q_{SNB}^H	0.1	10	16	3	29	33	0	24	7	31	32	0	0	0	1	1
	0.3	10	16	6	33	38	0	24	16	40	41	0	0	1	4	5
	1	9	16	18	44	48	0	23	40	63	63	0	0	0	3	3
	3	9	16	45	70	74	0	23	105	128	129	0	0	1	10	11
	10	10	16	112	137	143			OOM			0	0	0	22	22
Q_{SNB}^I	0.1	9	17	2	28	32	0	24	3	27	27	0	0	0	0	0
	0.3	10	16	3	29	34	0	24	6	30	31	0	0	1	1	2
	1	9	16	8	33	38	0	24	14	38	39	0	0	0	1	1
	3	10	16	14	39	43	0	24	33	57	57	0	0	1	4	5
	10	9	15	40	64	69			OOM			0	0	0	6	7

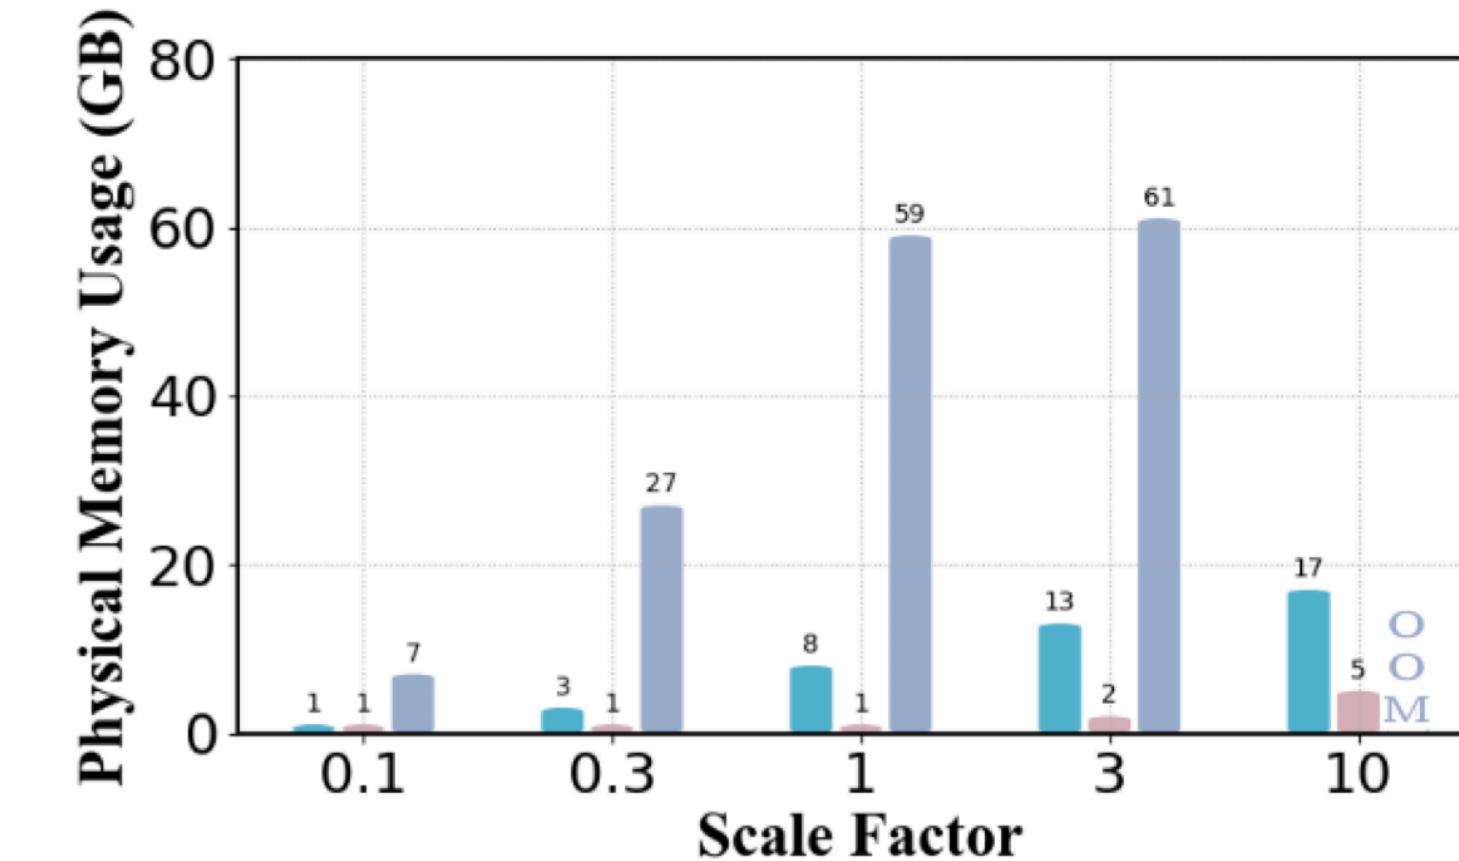
t_{Pa} = Parse time, t_{Pl} = Planning time, t_{Ex} = Execution time, t_{Bd} = Breakdown time, t_{tot} = End – to – end time

Results | LDBC SNB Benchmark

- **Physical memory efficiency:** Even at SF = 10, eStore peaks at 17 GB—nearly **4x lower** than Neo4jⁱ (61 GB) and well within system limits.
- **Graceful scaling:** eStore's physical usage **rises predictably** (1 → 17 GB) with data size, avoiding sudden spikes.
- **Take-away:** eStore delivers in-memory performance while **maintaining a far leaner** and more predictable memory profile than Neo4j's in-memory configuration.



(a) SNB virtual memory usage in GB.



(b) SNB physical memory usage in GB.

Results | Reimplementing Methods

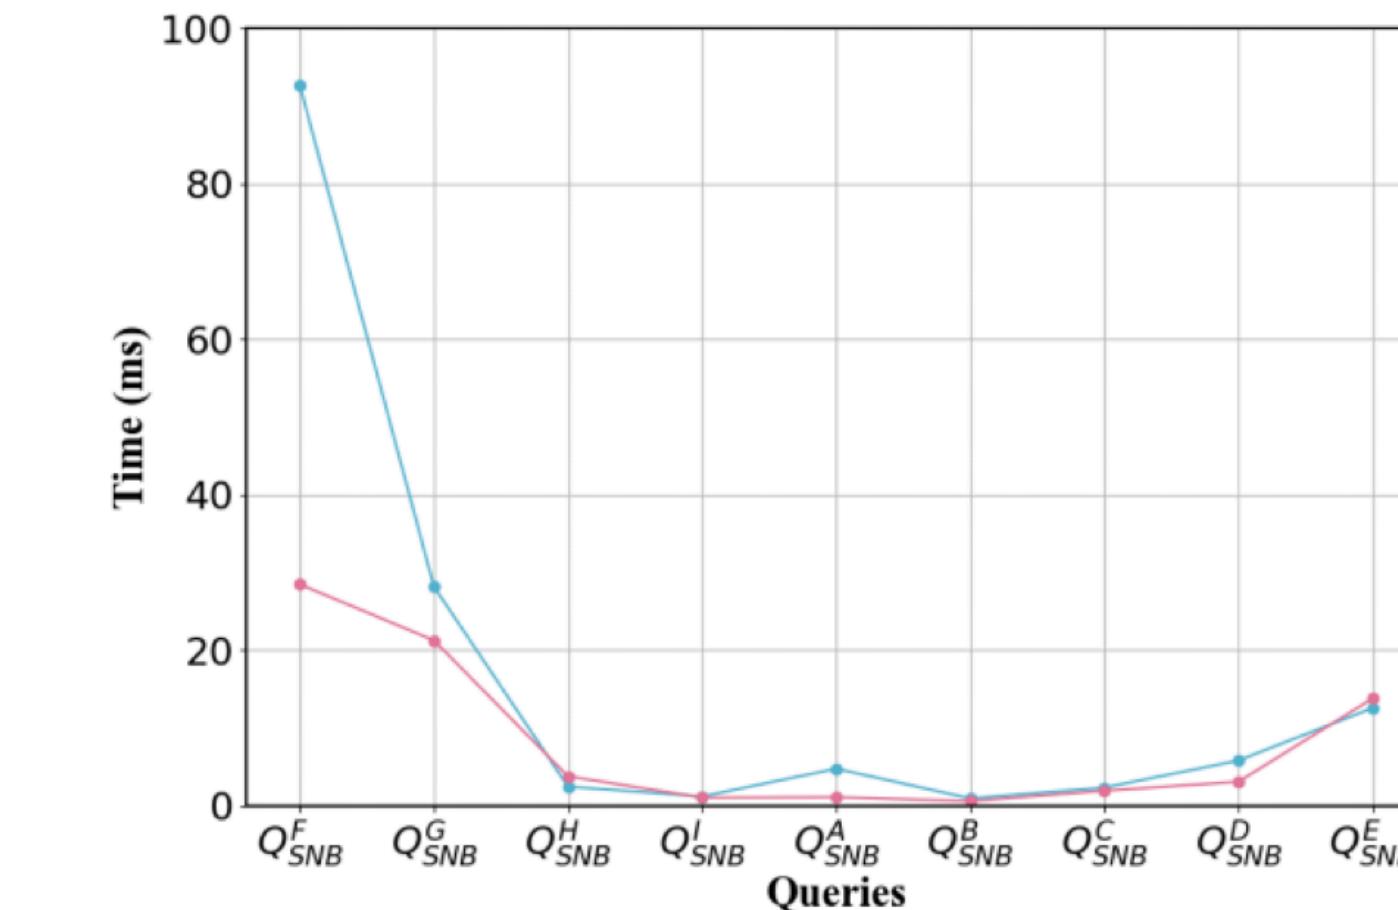
- **Benchmark setup:** measured end-to-end execution times for re-implementing data-structure methods in OGO vs. eStore.
- **Scalability:** OGO^{Neo} and OGO^{Mem} time out (> 1min) on larger workloads.
- **Performance:** eStore runs roughly **10x faster** than OGO across all workloads.

Data-structure	#Elements	OGO^{Neo}	OGO^{Mem}	ϵ_{Store}^r	ϵ_{Store}^u
JCF v17.0					
ArrayList	10^2	2231	264	89	81
	10^3	8169	486	97	90
	10^4	40261	888	133	116
	10^5	TO	TO	166	162
ArrayDeque	10^2	1632	399	88	86
	10^3	3377	497	105	89
	10^4	41301	1540	128	120
	10^5	TO	TO	156	166
HashMap	10^2	2391	315	78	77
	10^3	9479	516	88	100
	10^4	TO	2753	136	126
	10^5	TO	TO	356	394
LinkedList	10^2	24831	420	85	79
	10^3	TO	TO	83	88
	10^4	TO	TO	125	133
	10^5	TO	TO	232	216
Guava v32.1.3-jre					
ArrayList	10^2	1711	358	82	76
	10^3	2742	415	94	90
	10^4	23298	1031	125	140
	10^5	TO	TO	237	228
Eclipse v11.1.0					
UnifiedSet	10^2	1592	514	78	76
	10^3	2916	595	87	90
	10^4	55464	3373	158	149
	10^5	TO	TO	223	239
UnifiedMap	10^2	1997	489	88	83
	10^3	5079	865	102	109
	10^4	TO	TO	111	113
	10^5	TO	TO	210	200
FastList	10^2	1637	453	80	77
	10^3	2750	540	89	87
	10^4	40243	1255	130	143
	10^5	TO	TO	172	161
ArrayList	10^2	1820	411	78	82
	10^3	3538	741	87	84
	10^4	TO	2461	127	129
	10^5	TO	TO	153	153

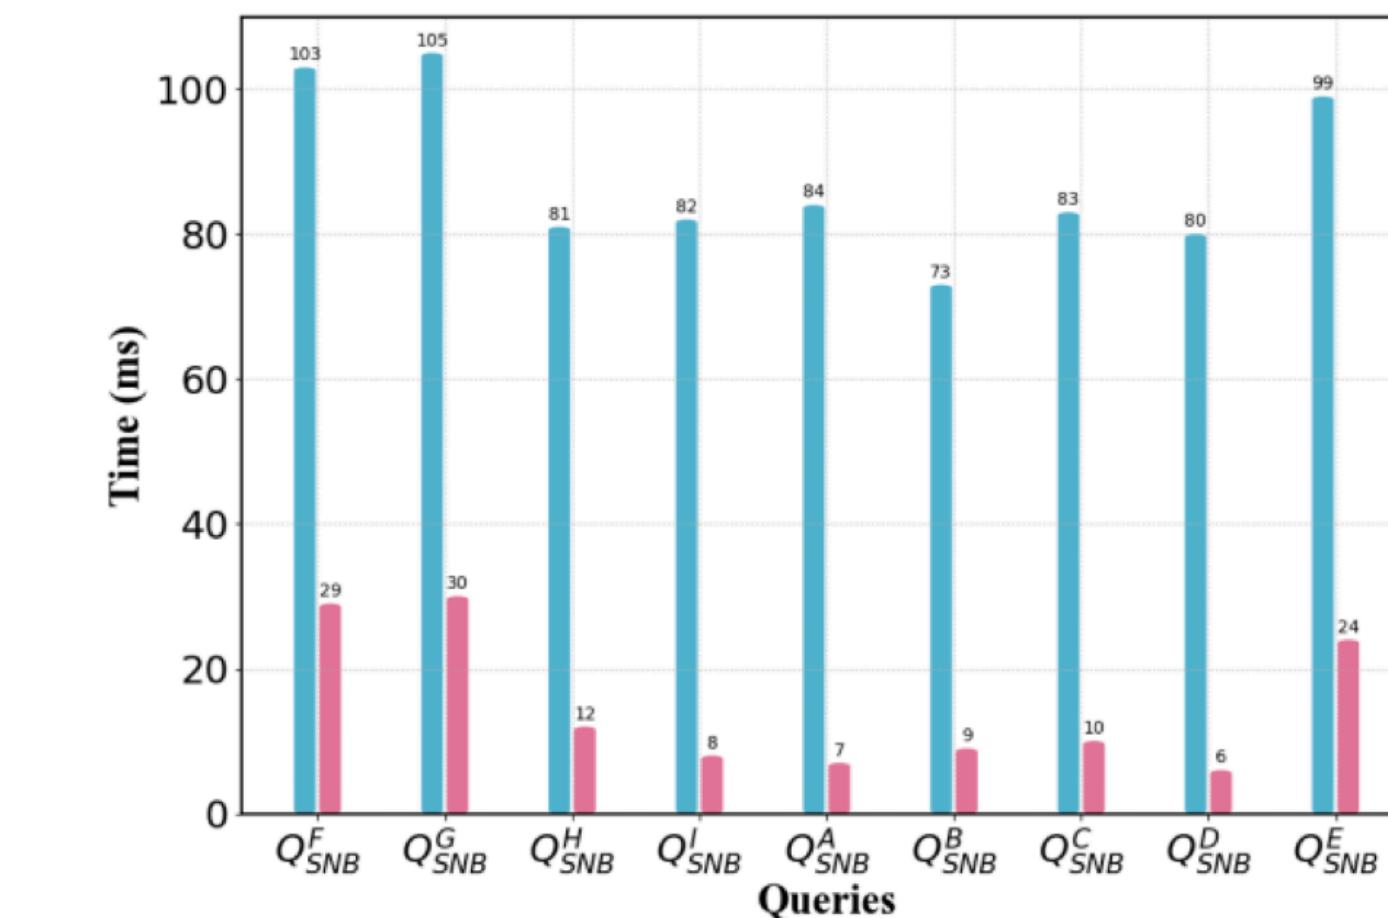
Results | Compile-time Code Generation

- **Benchmark setup:** LDBC SNB end-to-end times for eStore with vs. without compile-time rewriting.
- **Sequential run** (queries executed back-to-back): compile-time rewriting delivers a substantial speed-up on the first query.
- **Standalone run** (each query executed in isolation): compile-time rewriting accelerates every query.

■ ϵStore^r ■ ϵStore^r w/ Code Rewriting



(a) Sequence.



(b) Standalone.

Conclusions

- **Epsilon storage model:** queries live objects with virtually **no abstraction overhead**.
- **eStore:** graph-based object store that blends declarative queries with imperative code.
- **Compile-time query rewriting:** cuts **runtime latency** by an order of magnitude.
- **Benchmark results:** eStore surpasses production graph databases on LDBC SNB, making it a strong **testing-environment substitute**.
- **Robustness & versatility:** proven on LDBC SNB and by re-implementing methods in large-scale open-source projects.
- **Plus more:** comprehensive API design and formal foundations detailed in the paper.

Contact : auditt@utexas.edu

Supplementary